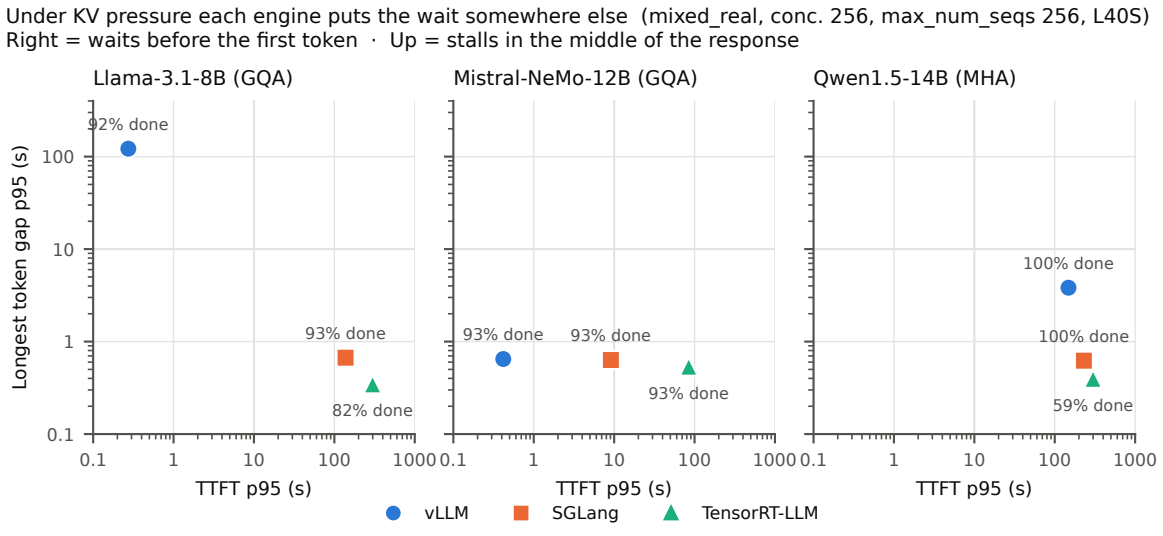
Under KV pressure each engine puts the wait somewhere else (mixed_real, conc. 256, max_num_seqs 256, L40S)
Right = waits before the first token · Up = stalls in the middle of the response
Llama-3.1-8B (GQA)
Mistral-NeMo-12B (GQA)
Qwen1.5-14B (MHA)
Longest token gap p95 (s)
TTFT p95 (s)
92% done
93% done
82% done
93% done
93% done
93% done
100% done
100% done
59% done
vLLM
SGLang
TensorRT-LLM