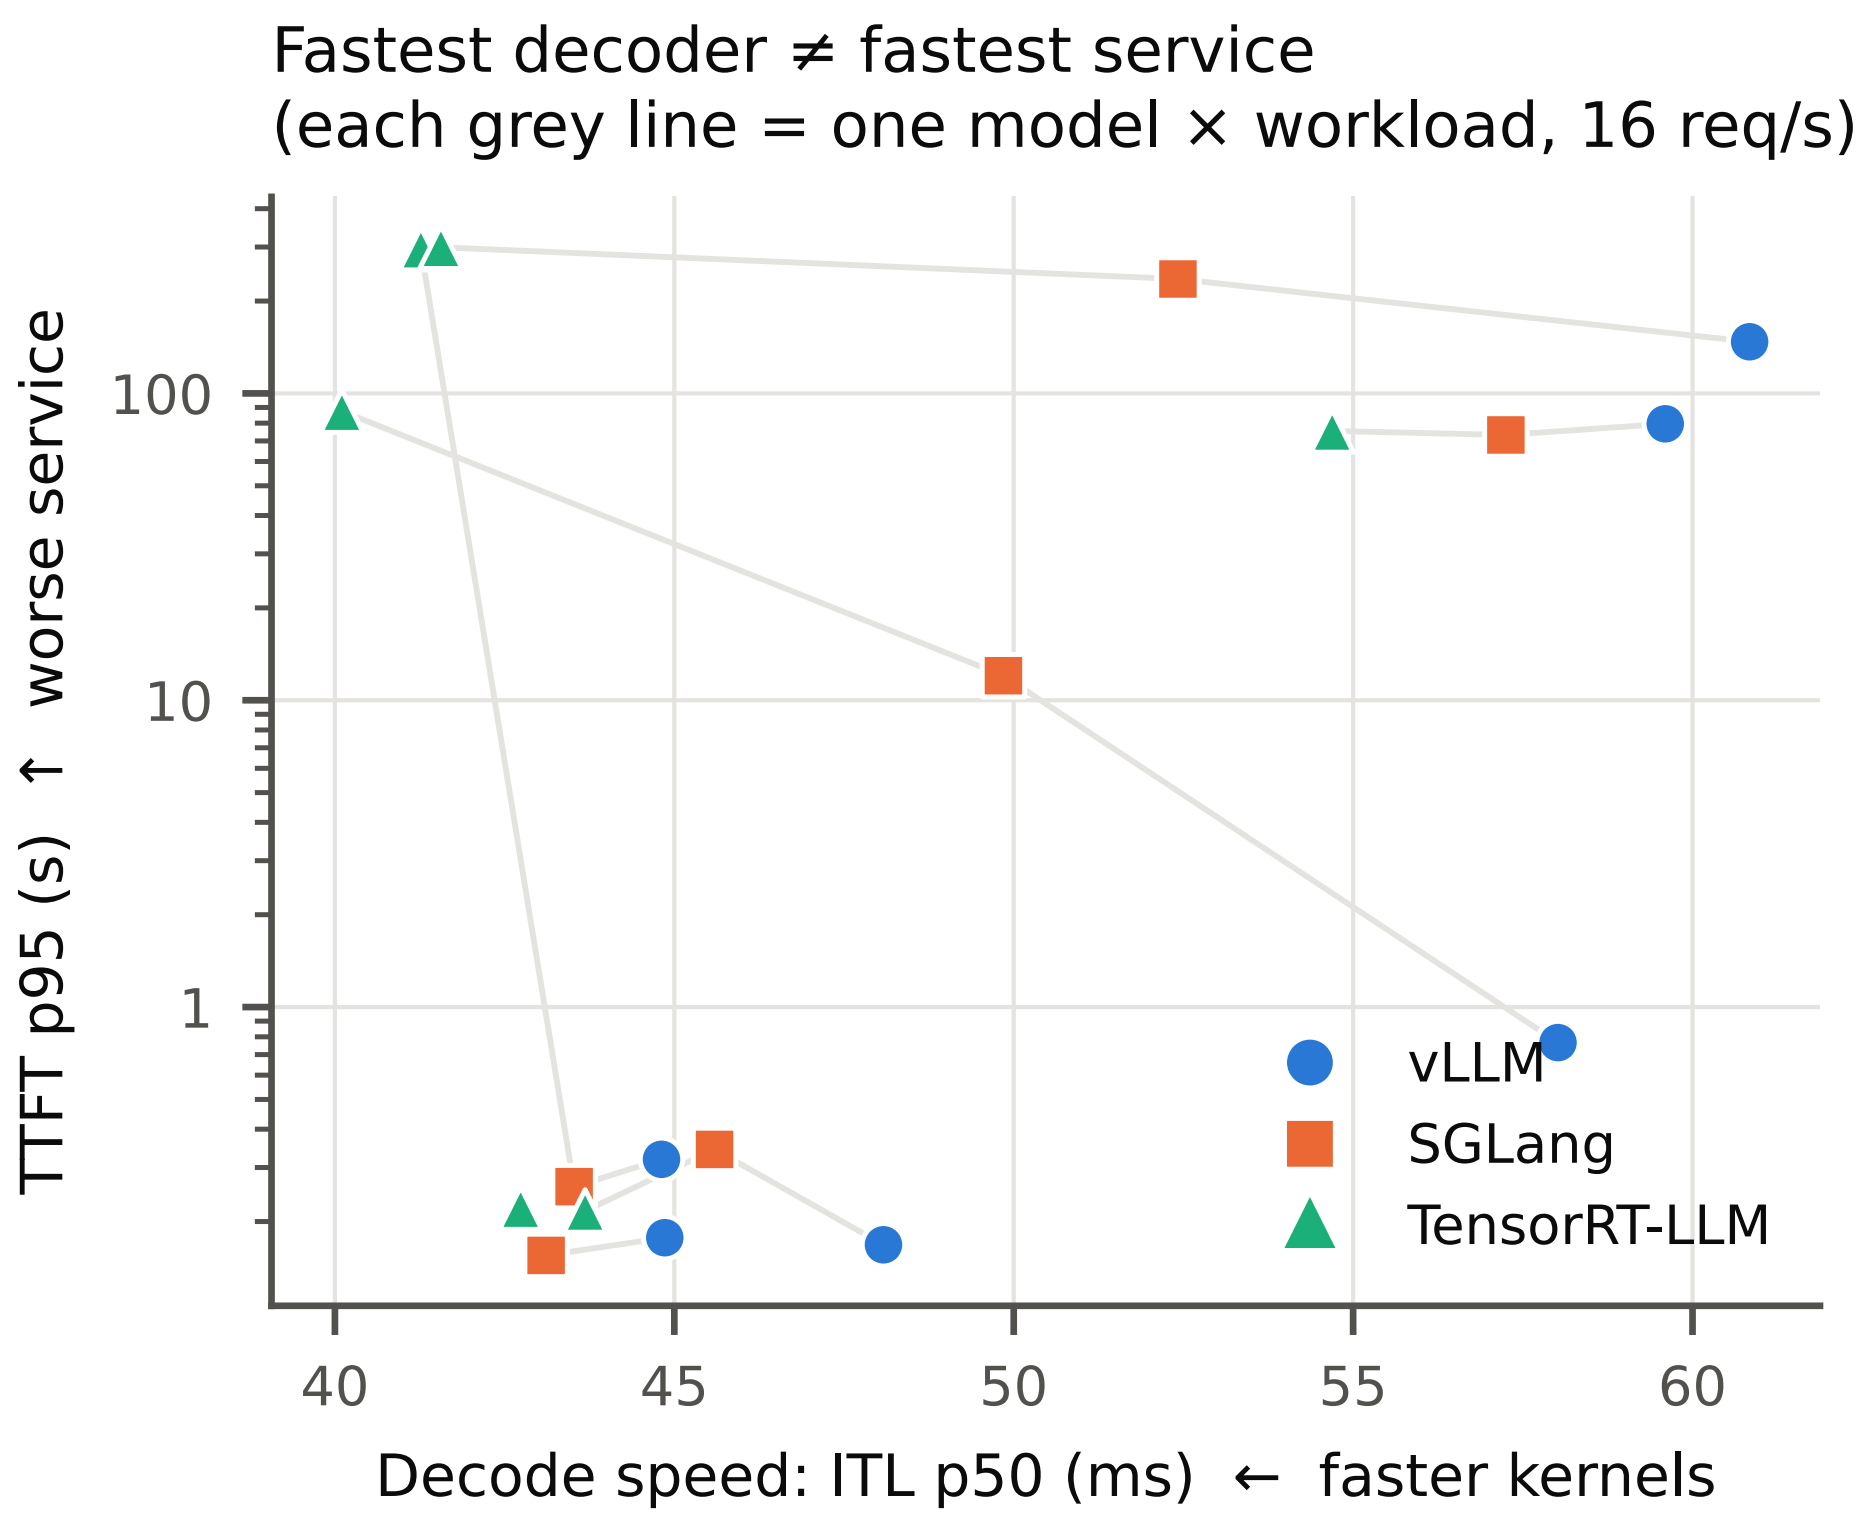
Fastest decoder ≠ fastest service
(each grey line = one model × workload, 16 req/s)
TTFT p95 (s) ↑ worse service
100
10
1
40
45
50
55
60
Decode speed: ITL p50 (ms) ← faster kernels
vLLM
SGLang
TensorRT-LLM