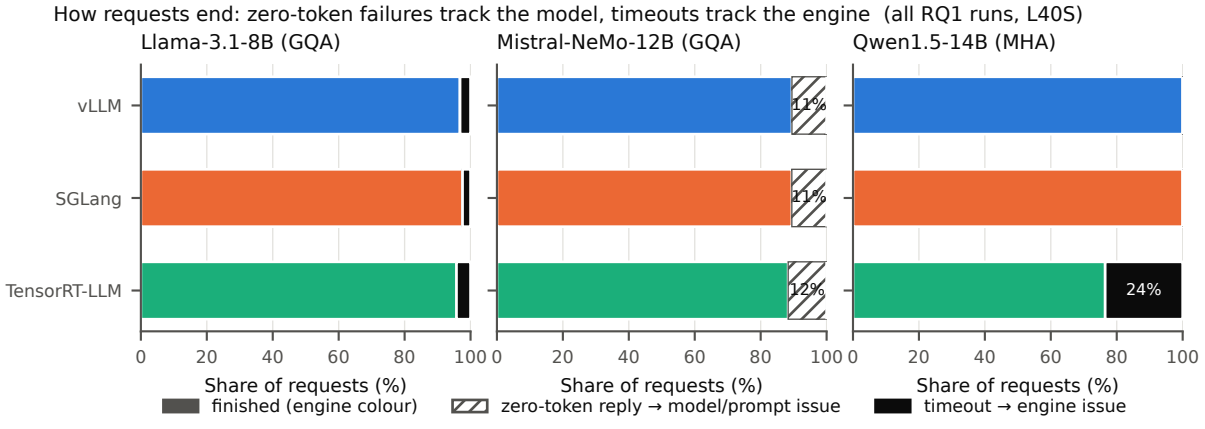
How requests end: zero-token failures track the model, timeouts track the engine (all RQ1 runs, L40S)
Llama-3.1-8B (GQA)
Mistral-NeMo-12B (GQA)
Qwen1.5-14B (MHA)
vLLM
SGLang
TensorRT-LLM
11%
11%
12%
24%
0
20
40
60
80
100
Share of requests (%)
finished (engine colour)
zero-token reply → model/prompt issue
timeout → engine issue