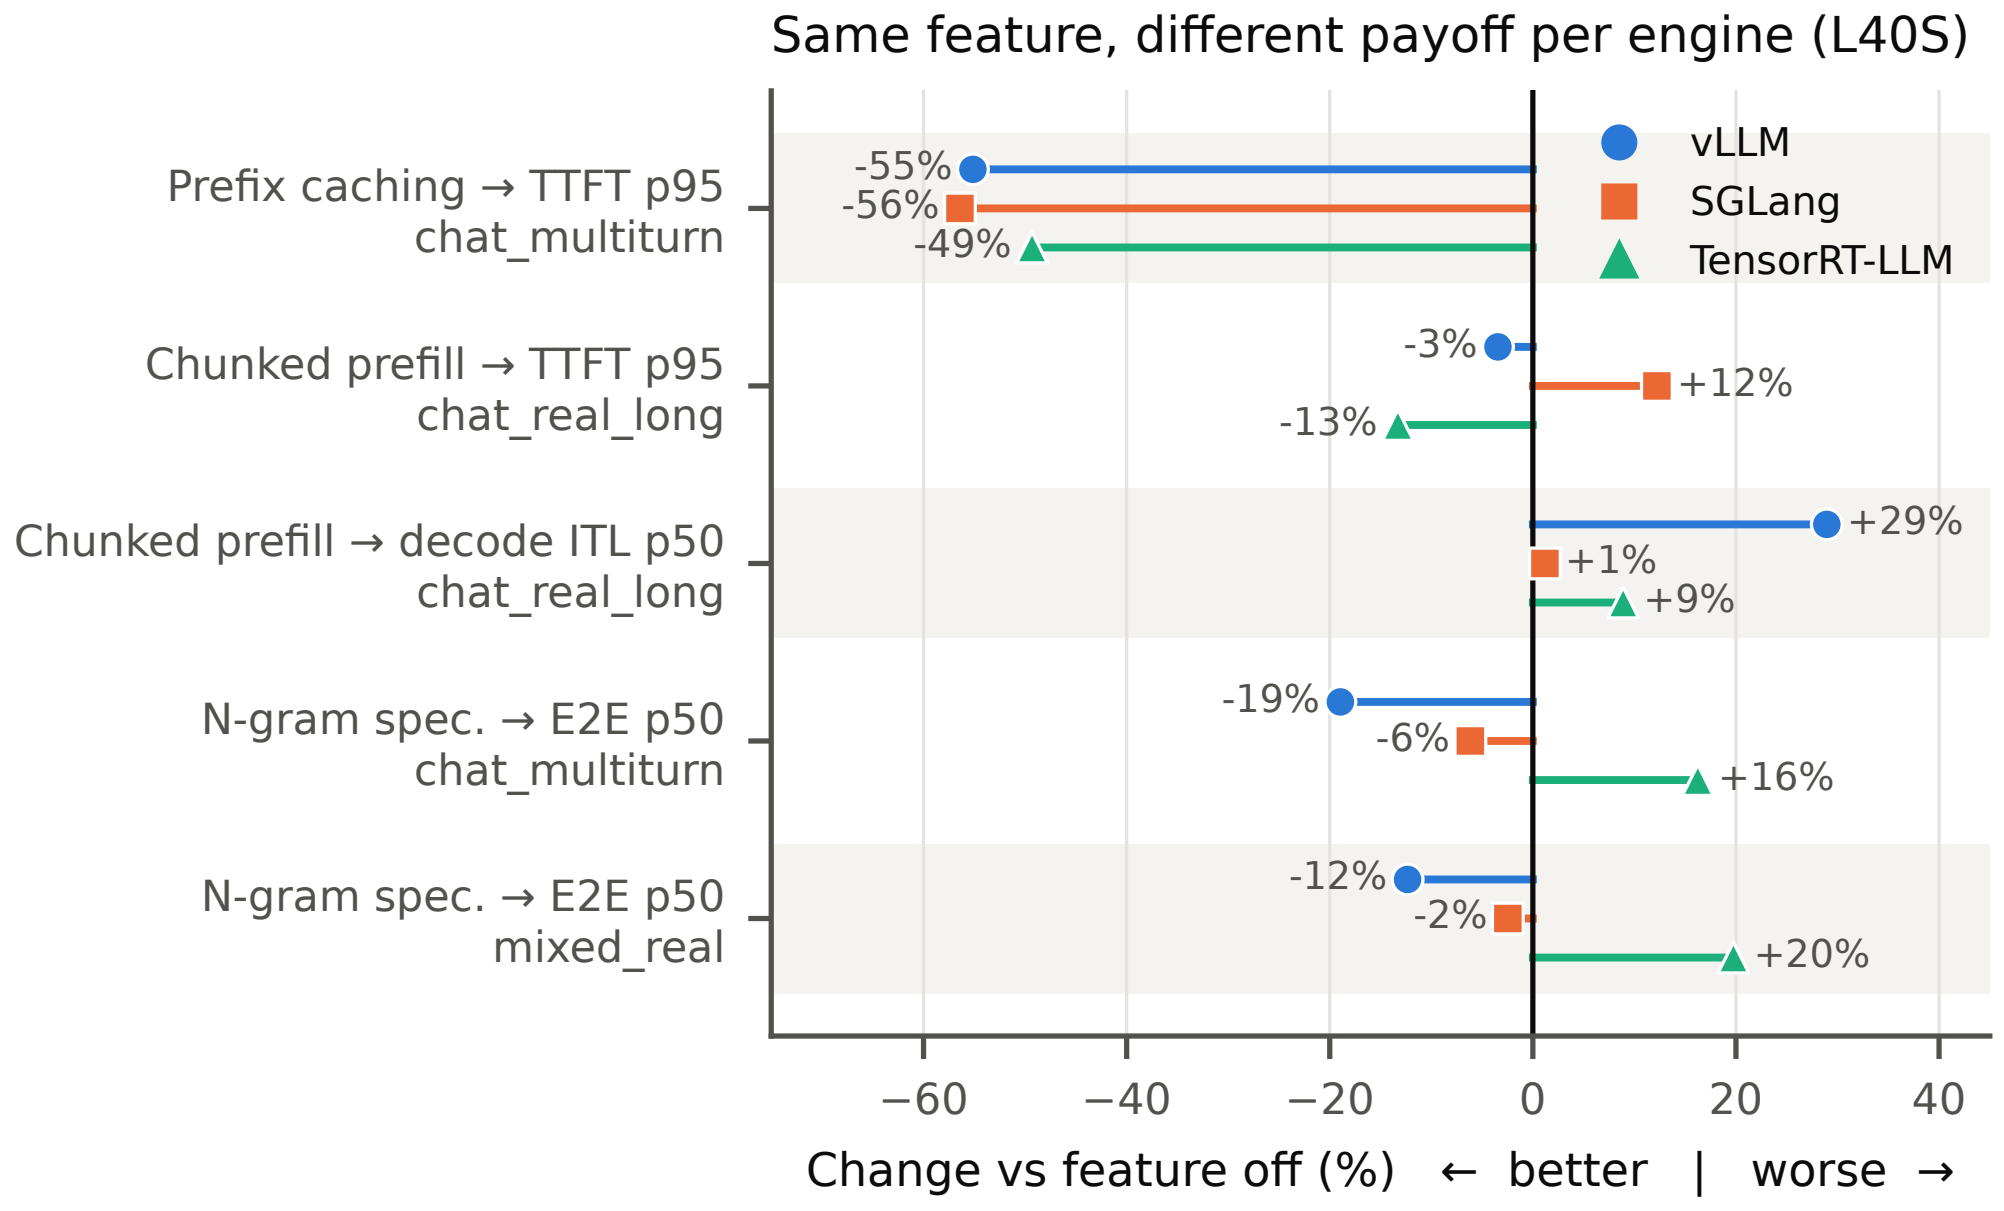
Same feature, different payoff per engine (L40S)
vLLM
SGLang
TensorRT-LLM
Prefix caching → TTFT p95
chat_multiturn
-55%
-56%
-49%
Chunked prefill → TTFT p95
chat_real_long
-3%
+12%
-13%
Chunked prefill → decode ITL p50
chat_real_long
+29%
+1%
+9%
N-gram spec. → E2E p50
chat_multiturn
-19%
-6%
+16%
N-gram spec. → E2E p50
mixed_real
-12%
-2%
+20%
−60
−40
−20
0
20
40
Change vs feature off (%) ← better | worse →