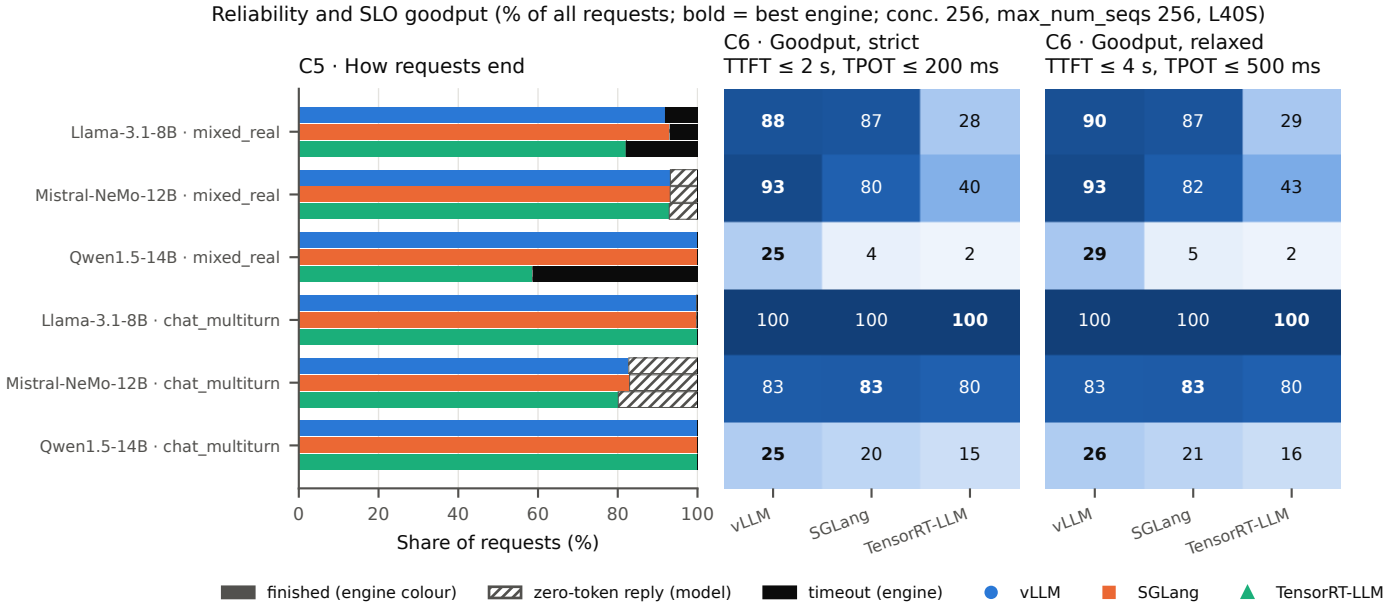
Reliability and SLO goodput (% of all requests; bold = best engine; conc. 256, max_num_seqs 256, L40S)
C5 · How requests end
C6 · Goodput, strict
TTFT ≤ 2 s, TPOT ≤ 200 ms
C6 · Goodput, relaxed
TTFT ≤ 4 s, TPOT ≤ 500 ms
Llama-3.1-8B · mixed_real
Mistral-NeMo-12B · mixed_real
Qwen1.5-14B · mixed_real
Llama-3.1-8B · chat_multiturn
Mistral-NeMo-12B · chat_multiturn
Qwen1.5-14B · chat_multiturn
0
20
40
60
80
100
Share of requests (%)
88 87 28
93 80 40
25 4 2
100 100 100
83 83 80
25 20 15
90 87 29
93 82 43
29 5 2
100 100 100
83 83 80
26 21 16
vLLM
SGLang
TensorRT-LLM
vLLM
SGLang
TensorRT-LLM
finished (engine colour)
zero-token reply (model)
timeout (engine)
vLLM
SGLang
TensorRT-LLM