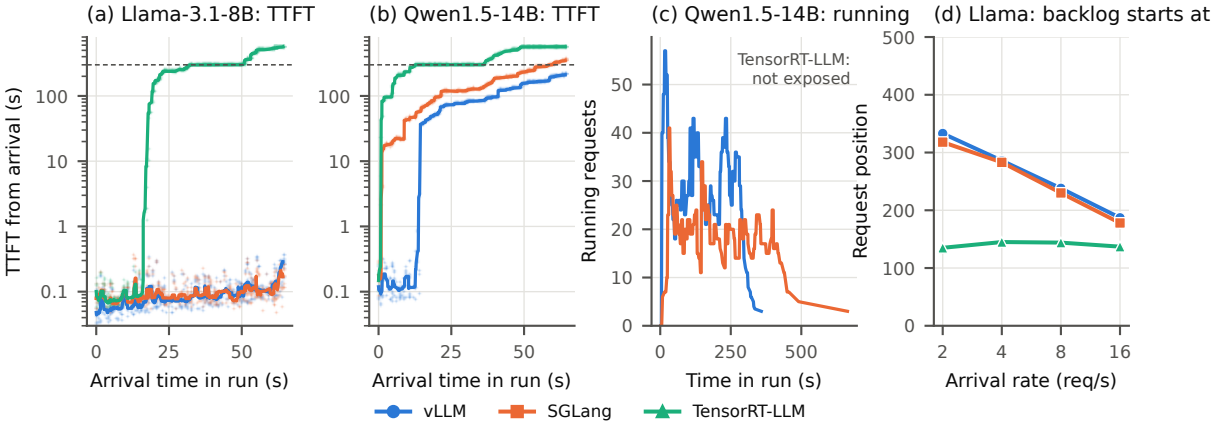
(a) Llama-3.1-8B: TTFT
(b) Qwen1.5-14B: TTFT
(c) Qwen1.5-14B: running
(d) Llama: backlog starts at
TTFT from arrival (s)
100
10
1
0.1
0
25
50
Arrival time in run (s)
Running requests
TensorRT-LLM:
not exposed
0
10
20
30
40
50
250
500
Time in run (s)
Request position
500
400
300
200
100
2
4
8
16
Arrival rate (req/s)
vLLM
SGLang
TensorRT-LLM