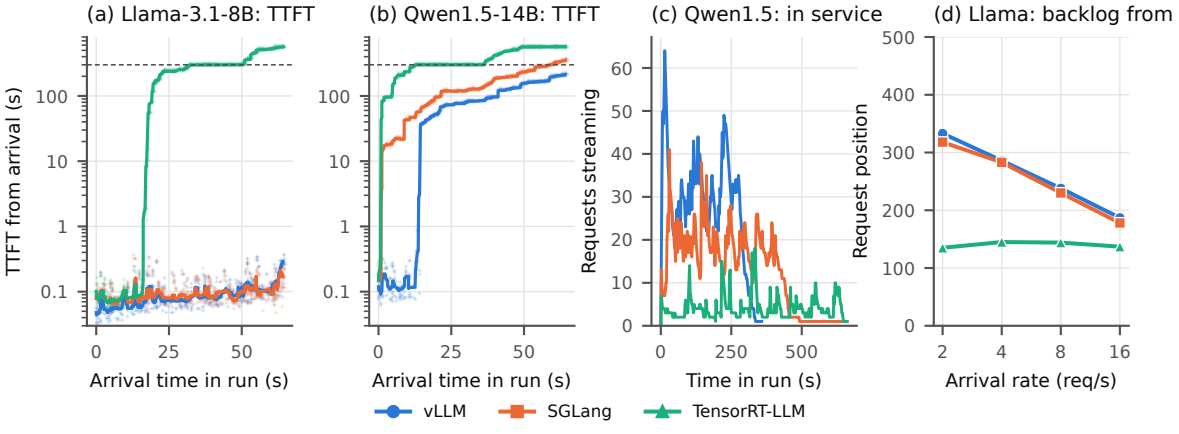
(a) Llama-3.1-8B: TTFT
(b) Qwen1.5-14B: TTFT
(c) Qwen1.5: in service
(d) Llama: backlog from
TTFT from arrival (s)
Requests streaming
Request position
Arrival time in run (s)
Time in run (s)
Arrival rate (req/s)
vLLM
SGLang
TensorRT-LLM