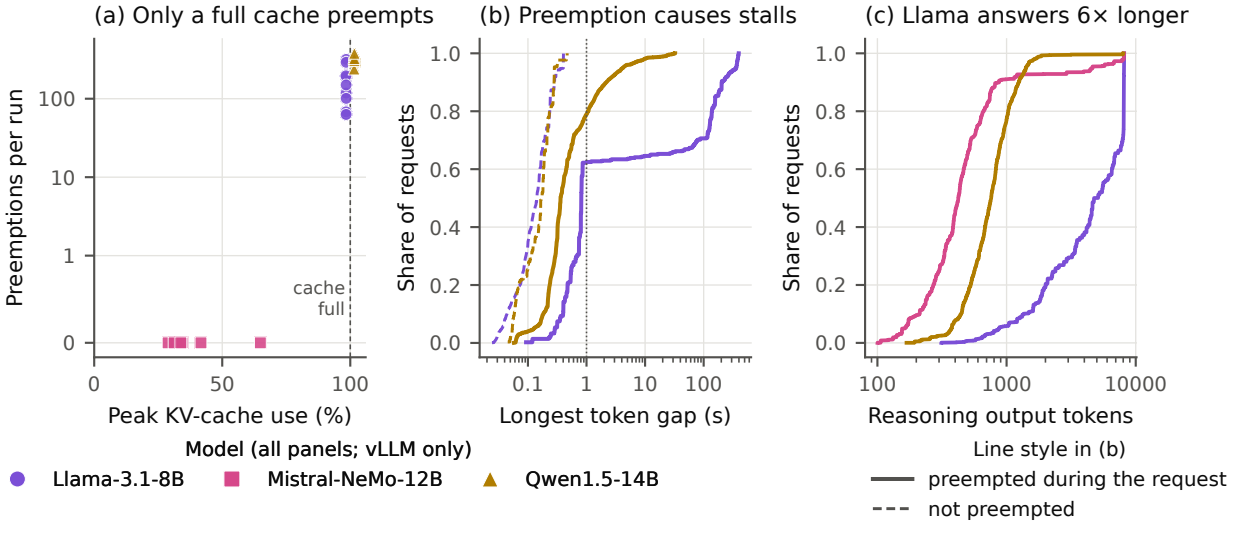
(a) Only a full cache preempts
(b) Preemption causes stalls
(c) Llama answers 6× longer
Preemptions per run
100
10
1
0
cache
full
0
50
100
Peak KV-cache use (%)
Share of requests
1.0
0.8
0.6
0.4
0.2
0.0
0.1
1
10
100
Longest token gap (s)
Share of requests
1.0
0.8
0.6
0.4
0.2
0.0
100
1000
10000
Reasoning output tokens
Model (all panels; vLLM only)
Llama-3.1-8B
Mistral-NeMo-12B
Qwen1.5-14B
Line style in (b)
preempted during the request
not preempted