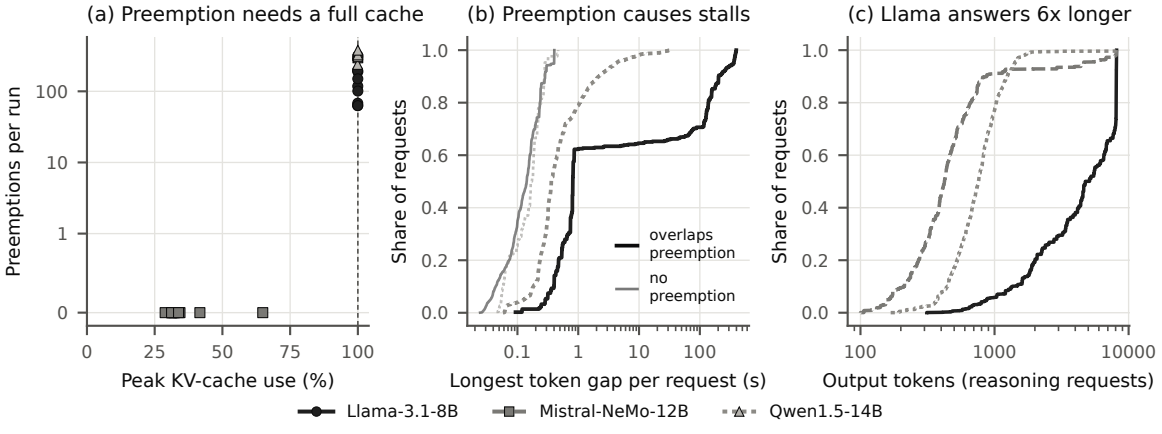
(a) Preemption needs a full cache
(b) Preemption causes stalls
(c) Llama answers 6x longer
Preemptions per run
Peak KV-cache use (%)
Share of requests
Longest token gap per request (s)
overlaps preemption
no preemption
Output tokens (reasoning requests)
Llama-3.1-8B
Mistral-NeMo-12B
Qwen1.5-14B