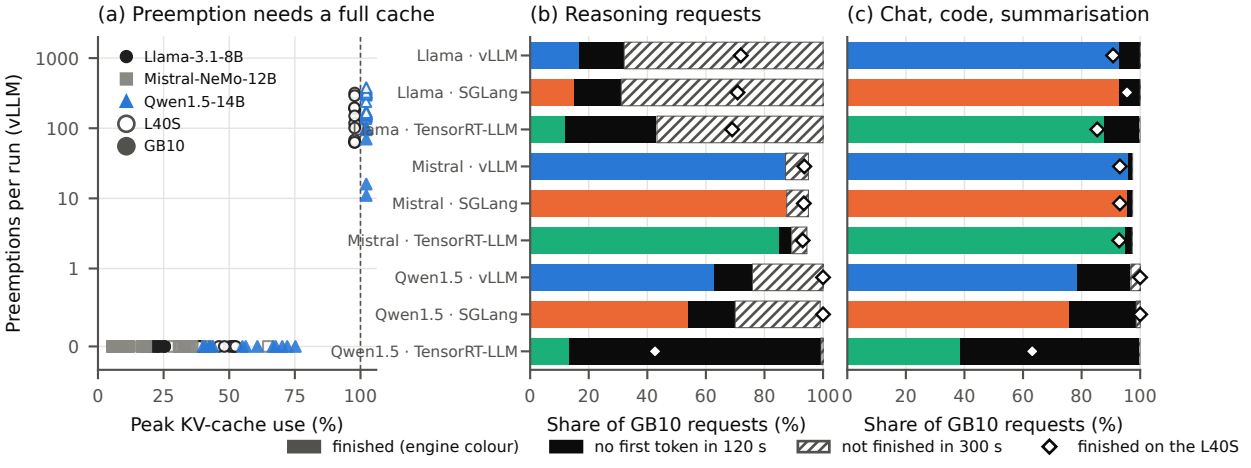
(a) Preemption needs a full cache
(b) Reasoning requests
(c) Chat, code, summarisation
Llama-3.1-8B
Mistral-NeMo-12B
Qwen1.5-14B
L40S
GB10
Preemptions per run (vLLM)
1000
100
10
1
0
0
25
50
75
100
Peak KV-cache use (%)
Llama · vLLM
Llama · SGLang
Llama · TensorRT-LLM
Mistral · vLLM
Mistral · SGLang
Mistral · TensorRT-LLM
Qwen1.5 · vLLM
Qwen1.5 · SGLang
Qwen1.5 · TensorRT-LLM
0
20
40
60
80
100
Share of GB10 requests (%)
Share of GB10 requests (%)
finished (engine colour)
no first token in 120 s
not finished in 300 s
finished on the L40S