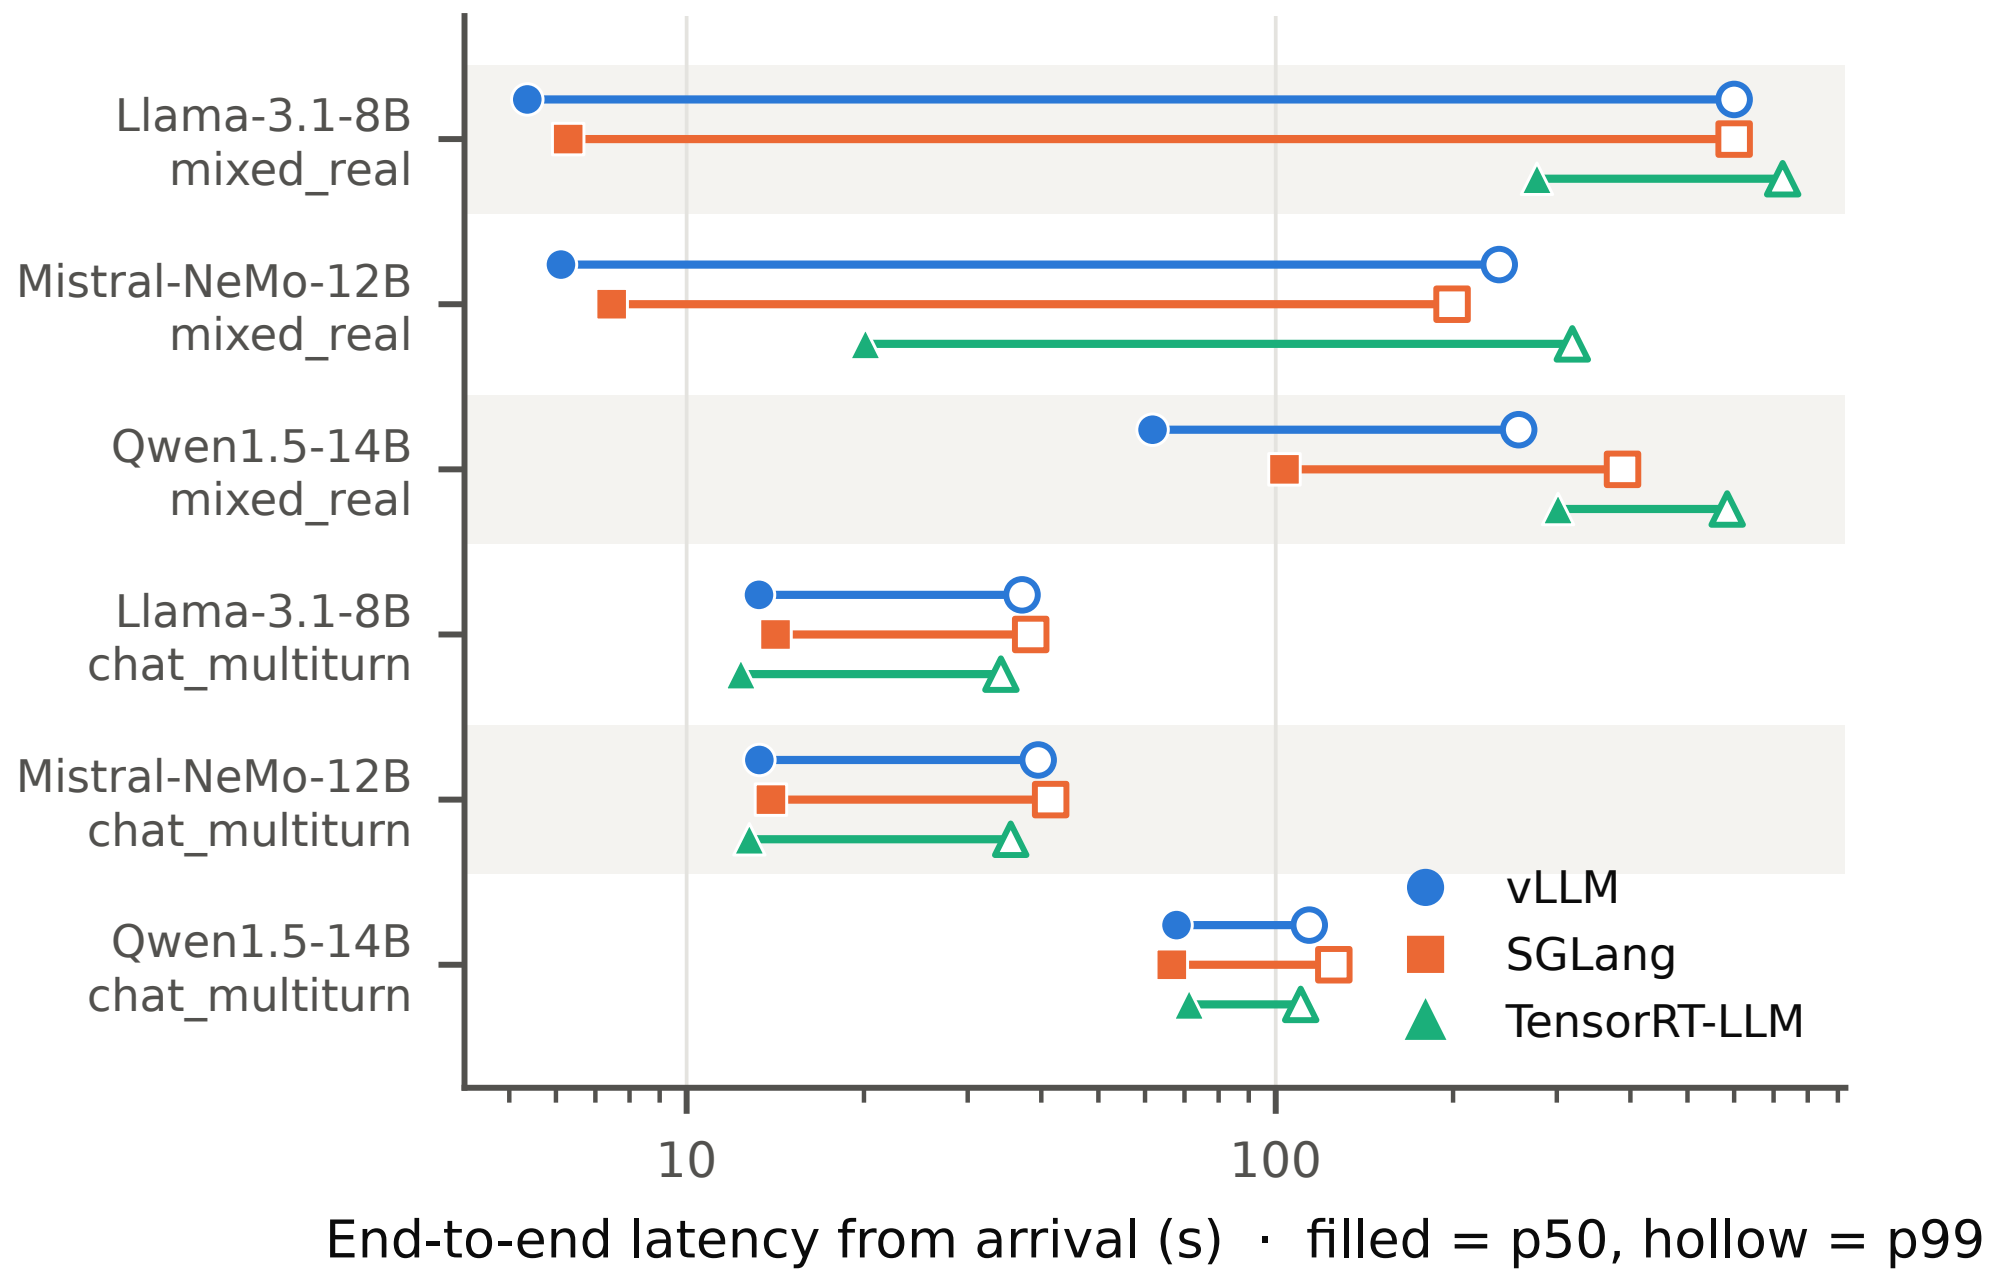
Llama-3.1-8B
mixed_real
Mistral-NeMo-12B
mixed_real
Qwen1.5-14B
mixed_real
Llama-3.1-8B
chat_multiturn
Mistral-NeMo-12B
chat_multiturn
Qwen1.5-14B
chat_multiturn
vLLM
SGLang
TensorRT-LLM
10
100
End-to-end latency from arrival (s) · filled = p50, hollow = p99