How requests end
Llama-3.1-8B · mixed_real
Mistral-NeMo-12B · mixed_real
Qwen1.5-14B · mixed_real
Llama-3.1-8B · chat_multiturn
Mistral-NeMo-12B · chat_multiturn
Qwen1.5-14B · chat_multiturn
0
20
40
60
80
100
Share of requests (%)
Strict goodput (%)
≤ 2 s TTFT, ≤ 200 ms TPOT
88 87 28
93 80 40
25 4 2
100 100 100
83 83 80
25 20 15
vLLM
SGLang
TensorRT-LLM
Relaxed goodput (%)
≤ 4 s TTFT, ≤ 500 ms TPOT
90 87 29
93 82 43
29 5 2
100 100 100
83 83 80
26 21 16
vLLM
SGLang
TensorRT-LLM
finished (engine colour)
zero-token reply (model)
timeout (engine)
vLLM
SGLang
TensorRT-LLM