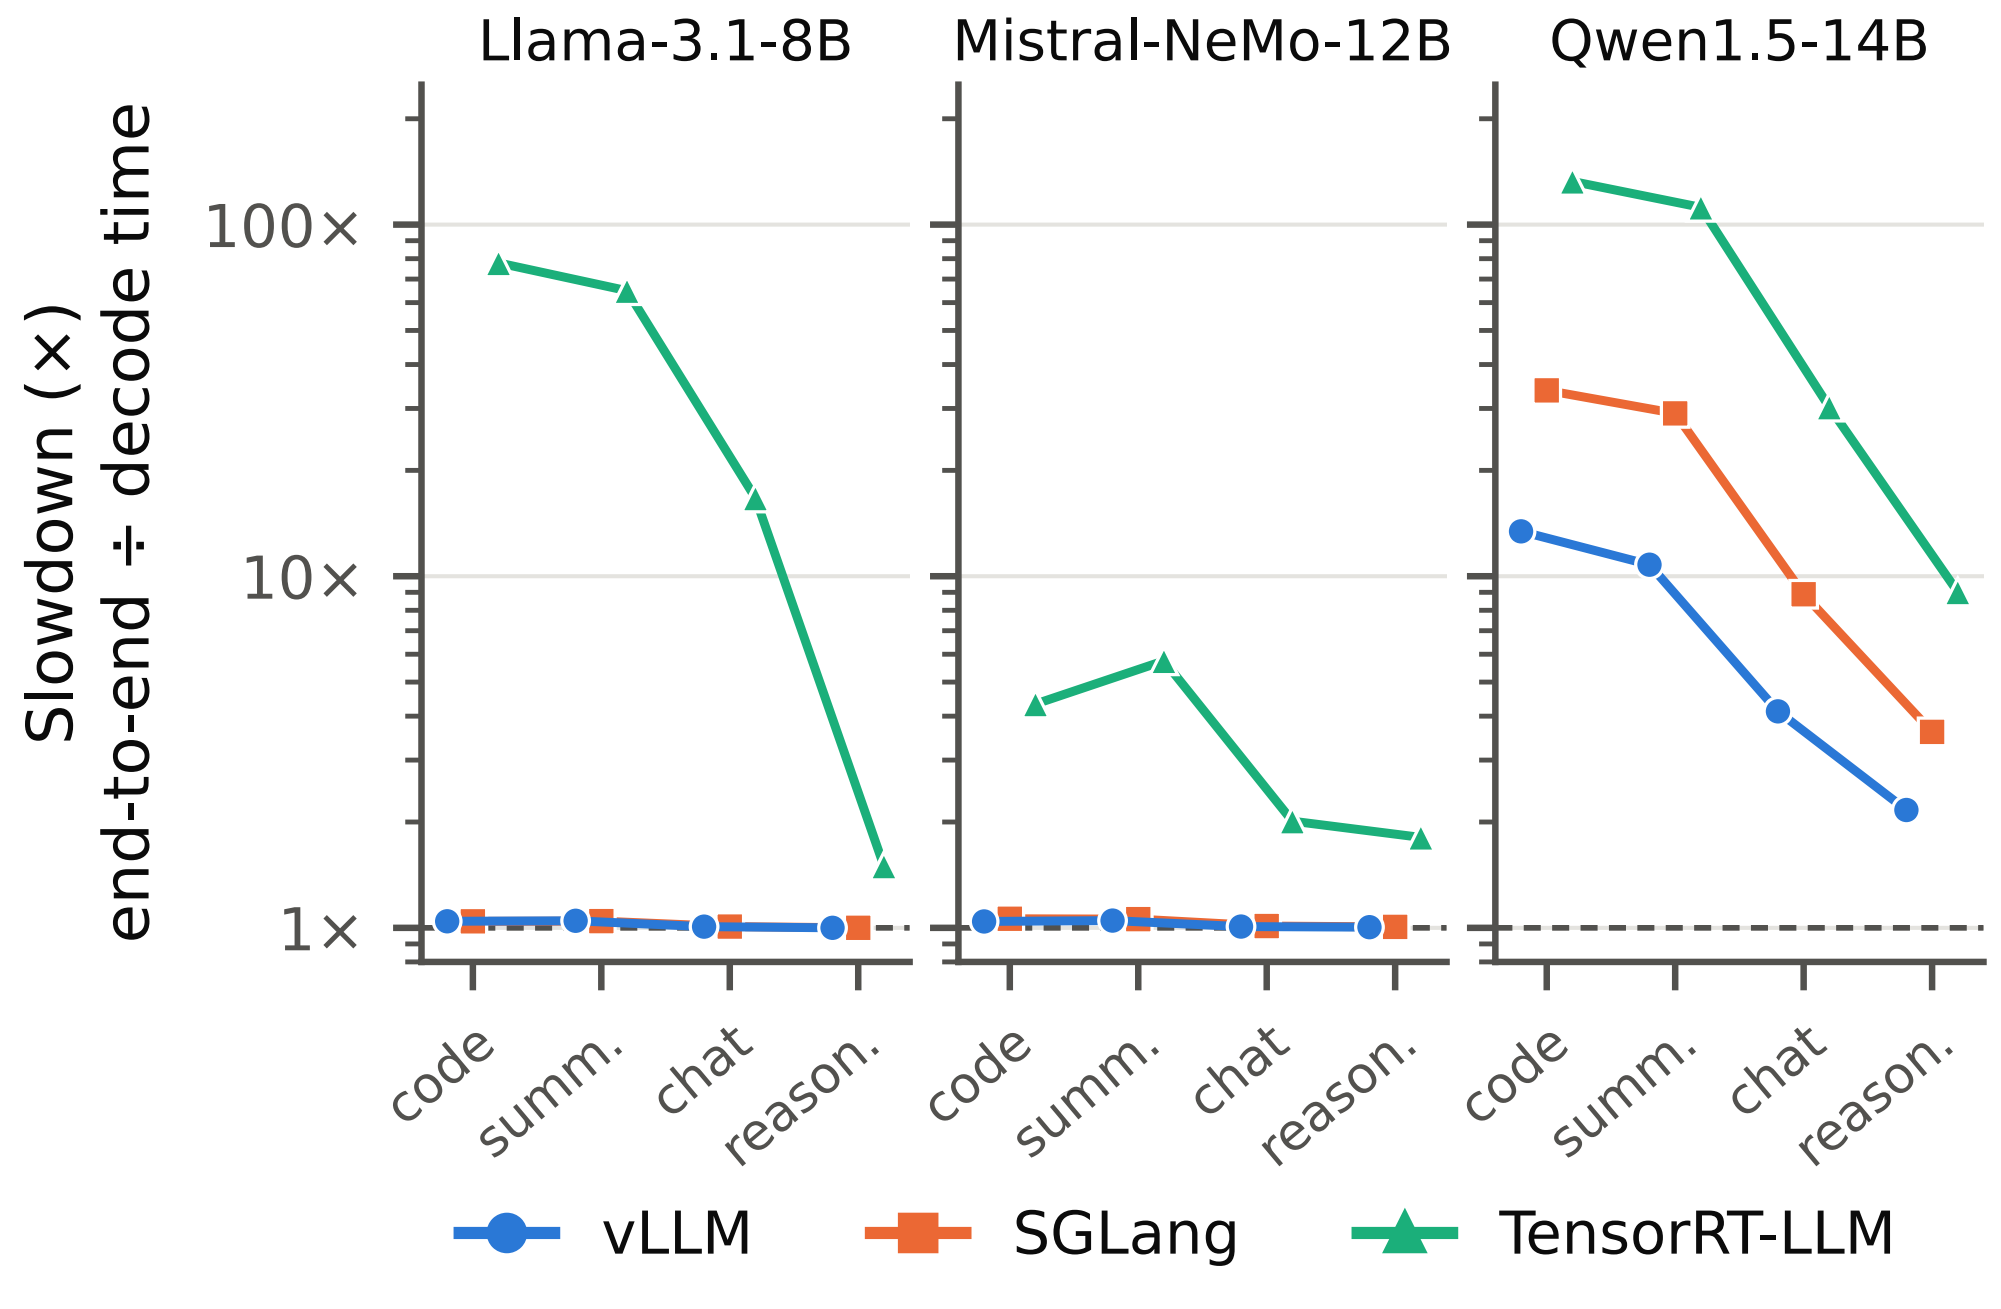
Llama-3.1-8B
Mistral-NeMo-12B
Qwen1.5-14B
Slowdown (×)
end-to-end ÷ decode time
100×
10×
1×
code
summ.
chat
reason.
vLLM
SGLang
TensorRT-LLM