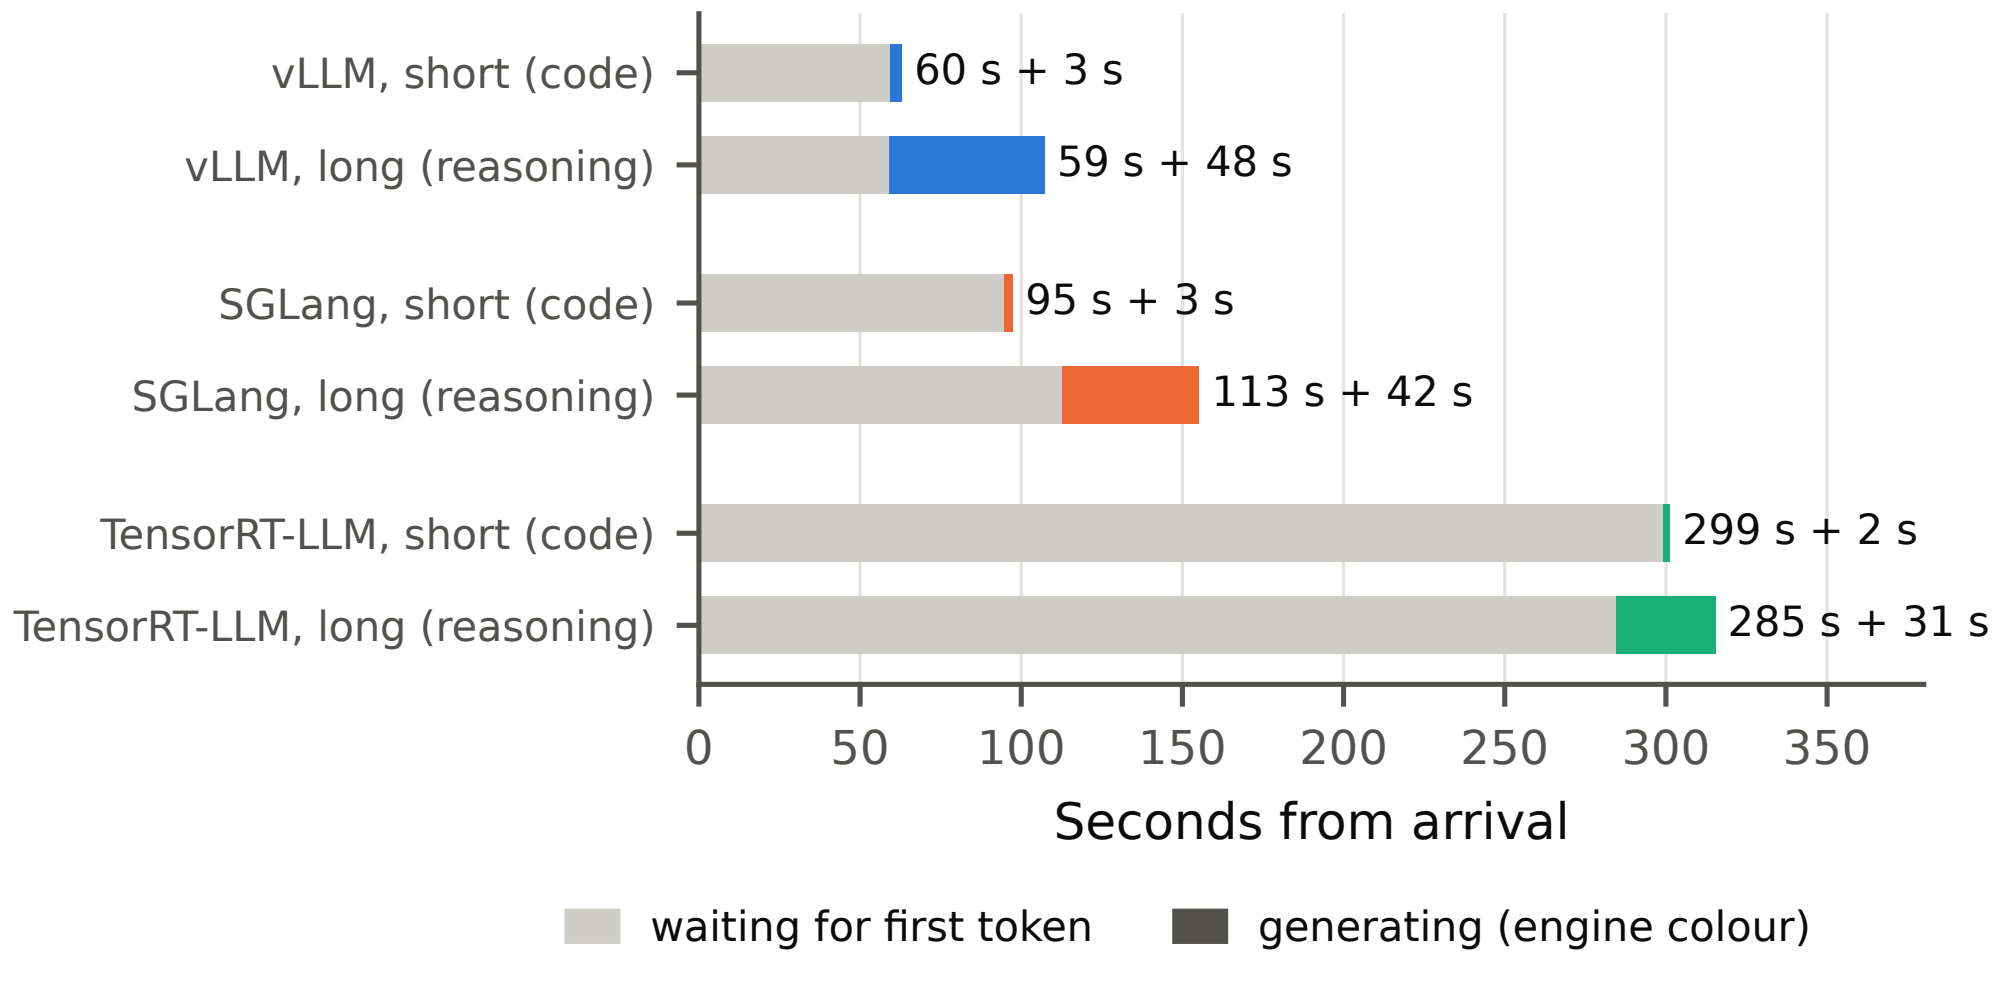
vLLM, short (code)
60 s + 3 s
vLLM, long (reasoning)
59 s + 48 s
SGLang, short (code)
95 s + 3 s
SGLang, long (reasoning)
113 s + 42 s
TensorRT-LLM, short (code)
299 s + 2 s
TensorRT-LLM, long (reasoning)
285 s + 31 s
0
50
100
150
200
250
300
350
Seconds from arrival
waiting for first token
generating (engine colour)