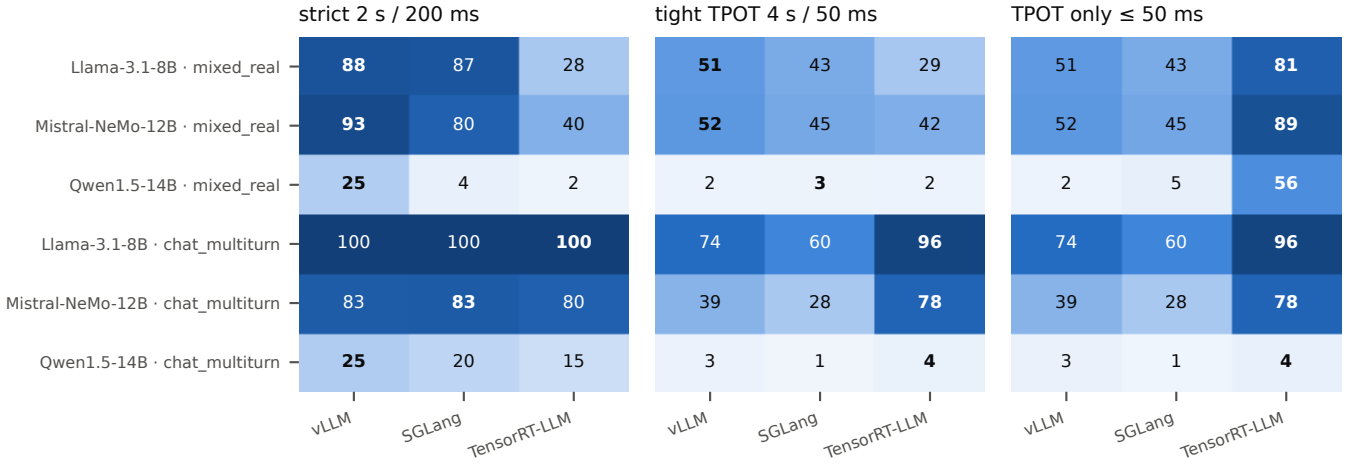
strict 2 s / 200 ms
tight TPOT 4 s / 50 ms
TPOT only ≤ 50 ms
Llama-3.1-8B · mixed_real
Mistral-NeMo-12B · mixed_real
Qwen1.5-14B · mixed_real
Llama-3.1-8B · chat_multiturn
Mistral-NeMo-12B · chat_multiturn
Qwen1.5-14B · chat_multiturn
88 87 28
93 80 40
25 4 2
100 100 100
83 83 80
25 20 15
51 43 29
52 45 42
2 3 2
74 60 96
39 28 78
3 1 4
51 43 81
52 45 89
2 5 56
74 60 96
39 28 78
3 1 4
vLLM
SGLang
TensorRT-LLM