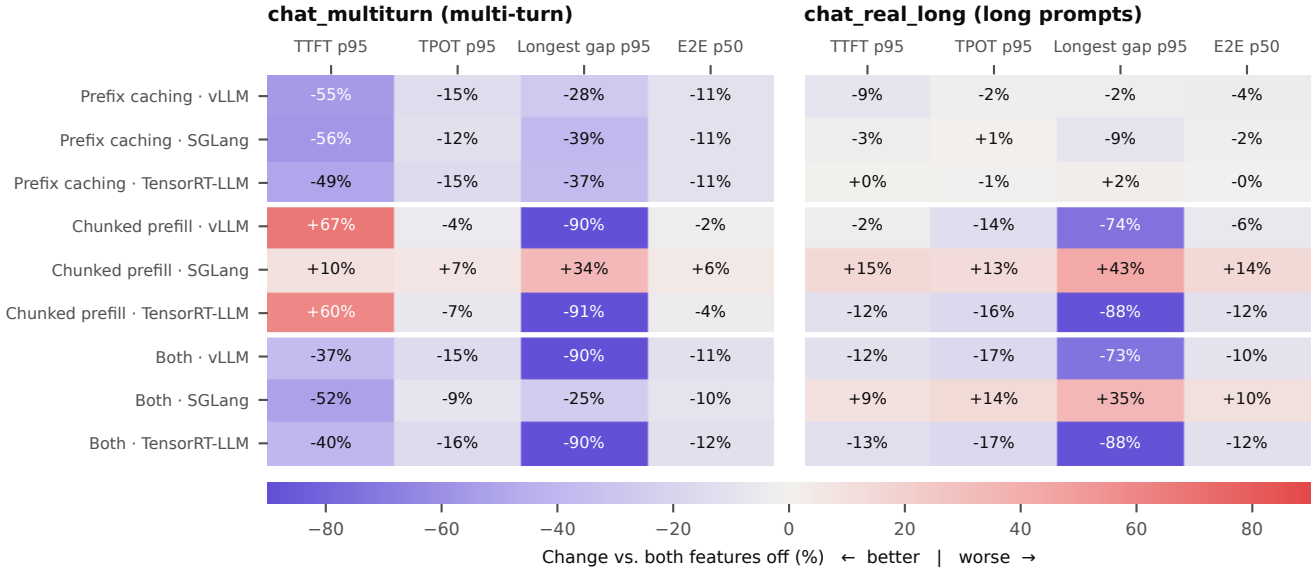
chat_multiturn (multi-turn)
TTFT p95
TPOT p95
Longest gap p95
E2E p50
chat_real_long (long prompts)
TTFT p95
TPOT p95
Longest gap p95
E2E p50
Prefix caching · vLLM
-55%
-15%
-28%
-11%
-9%
-2%
-2%
-4%
Prefix caching · SGLang
-56%
-12%
-39%
-11%
-3%
+1%
-9%
-2%
Prefix caching · TensorRT-LLM
-49%
-15%
-37%
-11%
+0%
-1%
+2%
-0%
Chunked prefill · vLLM
+67%
-4%
-90%
-2%
-2%
-14%
-74%
-6%
Chunked prefill · SGLang
+10%
+7%
+34%
+6%
+15%
+13%
+43%
+14%
Chunked prefill · TensorRT-LLM
+60%
-7%
-91%
-4%
-12%
-16%
-88%
-12%
Both · vLLM
-37%
-15%
-90%
-11%
-12%
-17%
-73%
-10%
Both · SGLang
-52%
-9%
-25%
-10%
+9%
+14%
+35%
+10%
Both · TensorRT-LLM
-40%
-16%
-90%
-12%
-13%
-17%
-88%
-12%
−80
−60
−40
−20
0
20
40
60
80
Change vs. both features off (%) ← better | worse →