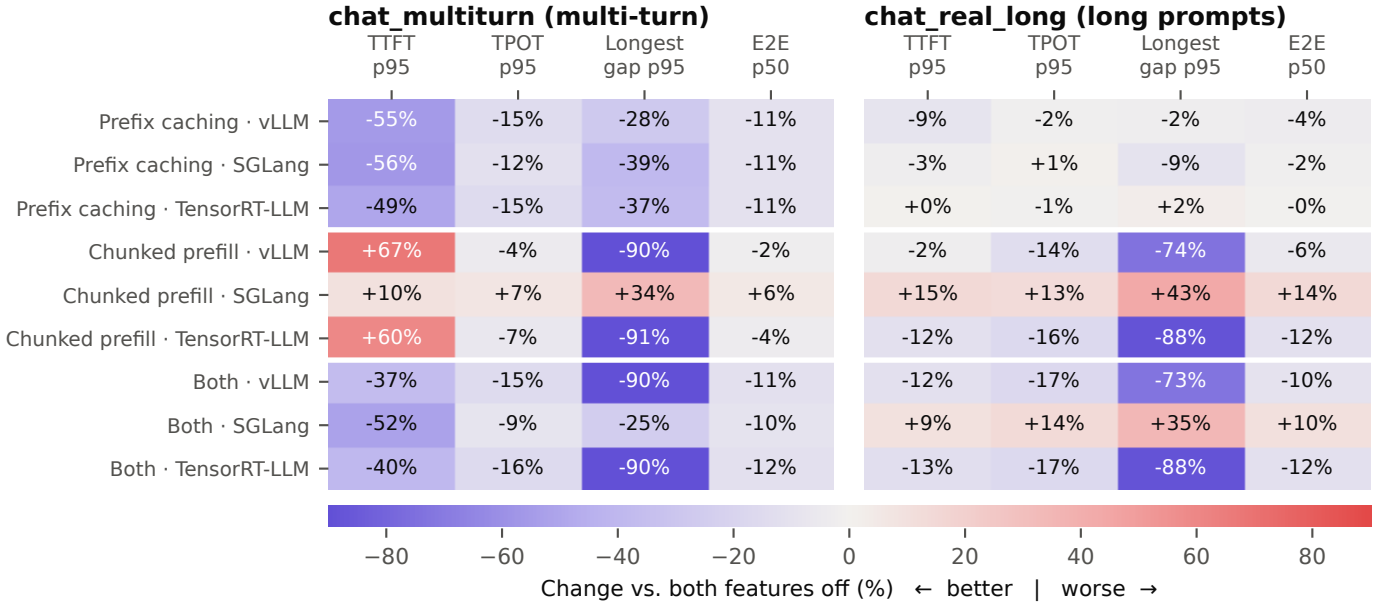
chat_multiturn (multi-turn)
chat_real_long (long prompts)
TTFT p95
TPOT p95
Longest gap p95
E2E p50
Prefix caching · vLLM
Prefix caching · SGLang
Prefix caching · TensorRT-LLM
Chunked prefill · vLLM
Chunked prefill · SGLang
Chunked prefill · TensorRT-LLM
Both · vLLM
Both · SGLang
Both · TensorRT-LLM
-55% -15% -28% -11% -9% -2% -2% -4%
-56% -12% -39% -11% -3% +1% -9% -2%
-49% -15% -37% -11% +0% -1% +2% -0%
+67% -4% -90% -2% -2% -14% -74% -6%
+10% +7% +34% +6% +15% +13% +43% +14%
+60% -7% -91% -4% -12% -16% -88% -12%
-37% -15% -90% -11% -12% -17% -73% -10%
-52% -9% -25% -10% +9% +14% +35% +10%
-40% -16% -90% -12% -13% -17% -88% -12%
−80 −60 −40 −20 0 20 40 60 80
Change vs. both features off (%) ← better | worse →