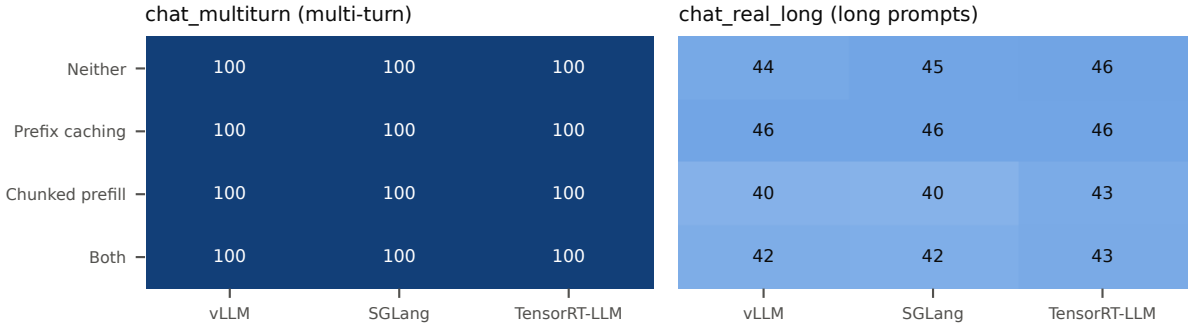
chat_multiturn (multi-turn)
chat_real_long (long prompts)
Neither
Prefix caching
Chunked prefill
Both
100
100
100
100
100
100
100
100
100
100
100
100
44
45
46
46
46
46
40
40
43
42
42
43
vLLM
SGLang
TensorRT-LLM
vLLM
SGLang
TensorRT-LLM