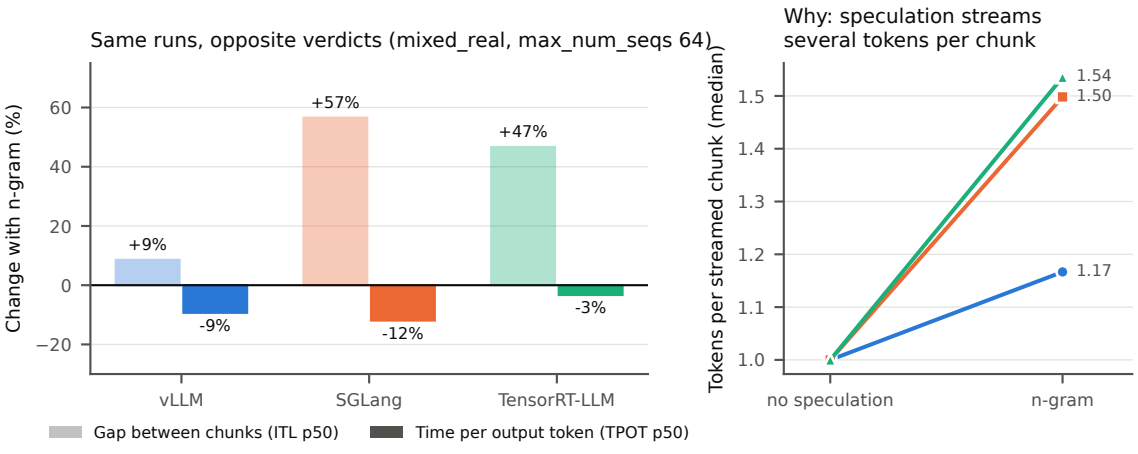
Same runs, opposite verdicts (mixed_real, max_num_seqs 64)
Change with n-gram (%)
60
40
20
0
−20
+9%
-9%
+57%
-12%
+47%
-3%
vLLM
SGLang
TensorRT-LLM
Gap between chunks (ITL p50)
Time per output token (TPOT p50)
Why: speculation streams
several tokens per chunk
Tokens per streamed chunk (median)
1.5
1.4
1.3
1.2
1.1
1.0
1.54
1.50
1.17
no speculation
n-gram