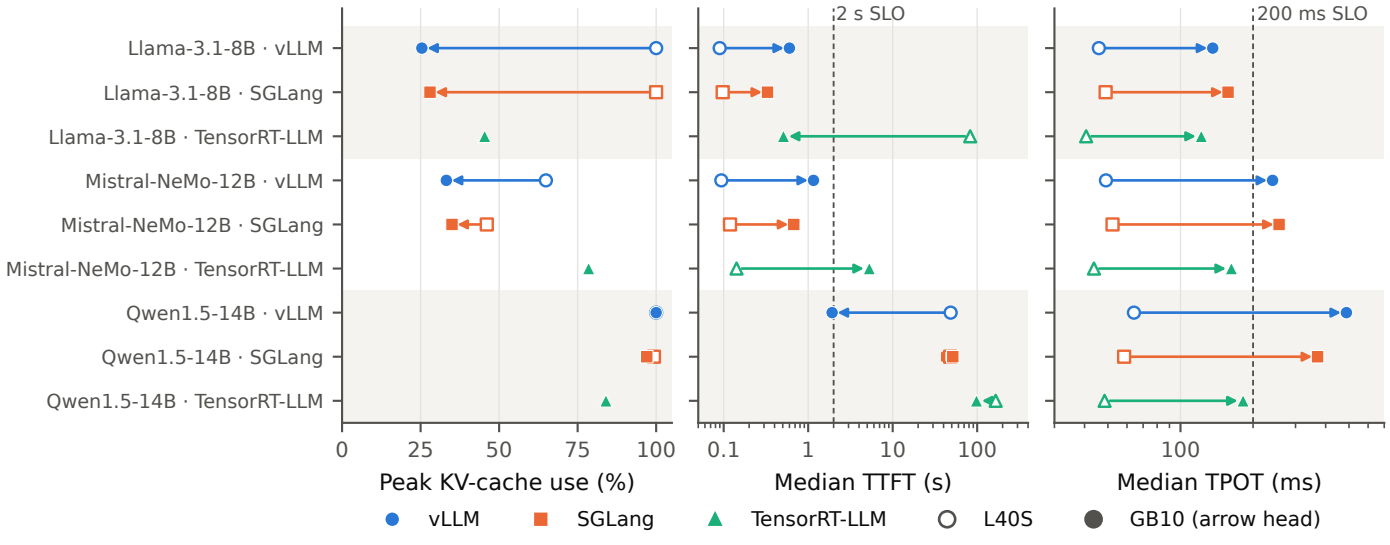
Llama-3.1-8B · vLLM
Llama-3.1-8B · SGLang
Llama-3.1-8B · TensorRT-LLM
Mistral-NeMo-12B · vLLM
Mistral-NeMo-12B · SGLang
Mistral-NeMo-12B · TensorRT-LLM
Qwen1.5-14B · vLLM
Qwen1.5-14B · SGLang
Qwen1.5-14B · TensorRT-LLM
2 s SLO
200 ms SLO
0
25
50
75
100
0.1
1
10
100
100
Peak KV-cache use (%)
Median TTFT (s)
Median TPOT (ms)
vLLM
SGLang
TensorRT-LLM
L40S
GB10 (arrow head)