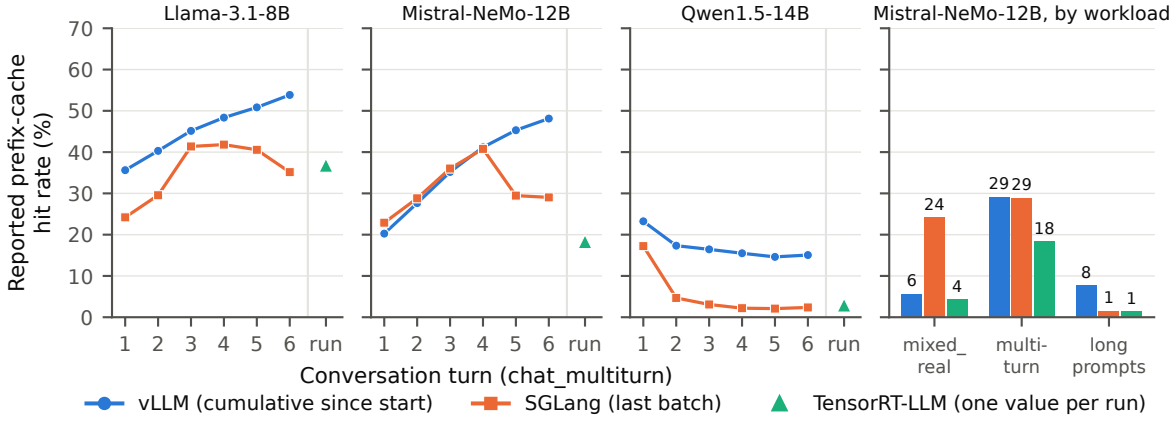
Llama-3.1-8B
Mistral-NeMo-12B
Qwen1.5-14B
Mistral-NeMo-12B, by workload
Reported prefix-cache hit rate (%)
0
10
20
30
40
50
60
70
1
2
3
4
5
6
run
Conversation turn (chat_multiturn)
6
24
4
29
29
18
8
1
1
mixed_ real
multi- turn
long prompts
vLLM (cumulative since start)
SGLang (last batch)
TensorRT-LLM (one value per run)