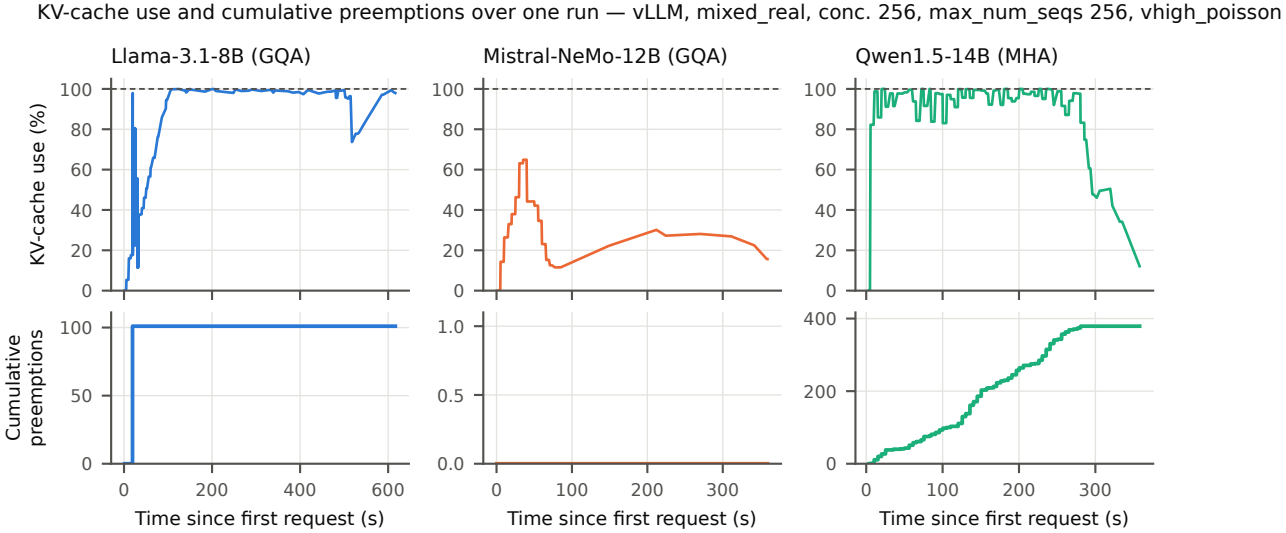
KV-cache use and cumulative preemptions over one run — vLLM, mixed_real, conc. 256, max_num_seqs 256, vhigh_poisson
Llama-3.1-8B (GQA)
Mistral-NeMo-12B (GQA)
Qwen1.5-14B (MHA)
KV-cache use (%)
Cumulative preemptions
Time since first request (s)
Time since first request (s)
Time since first request (s)