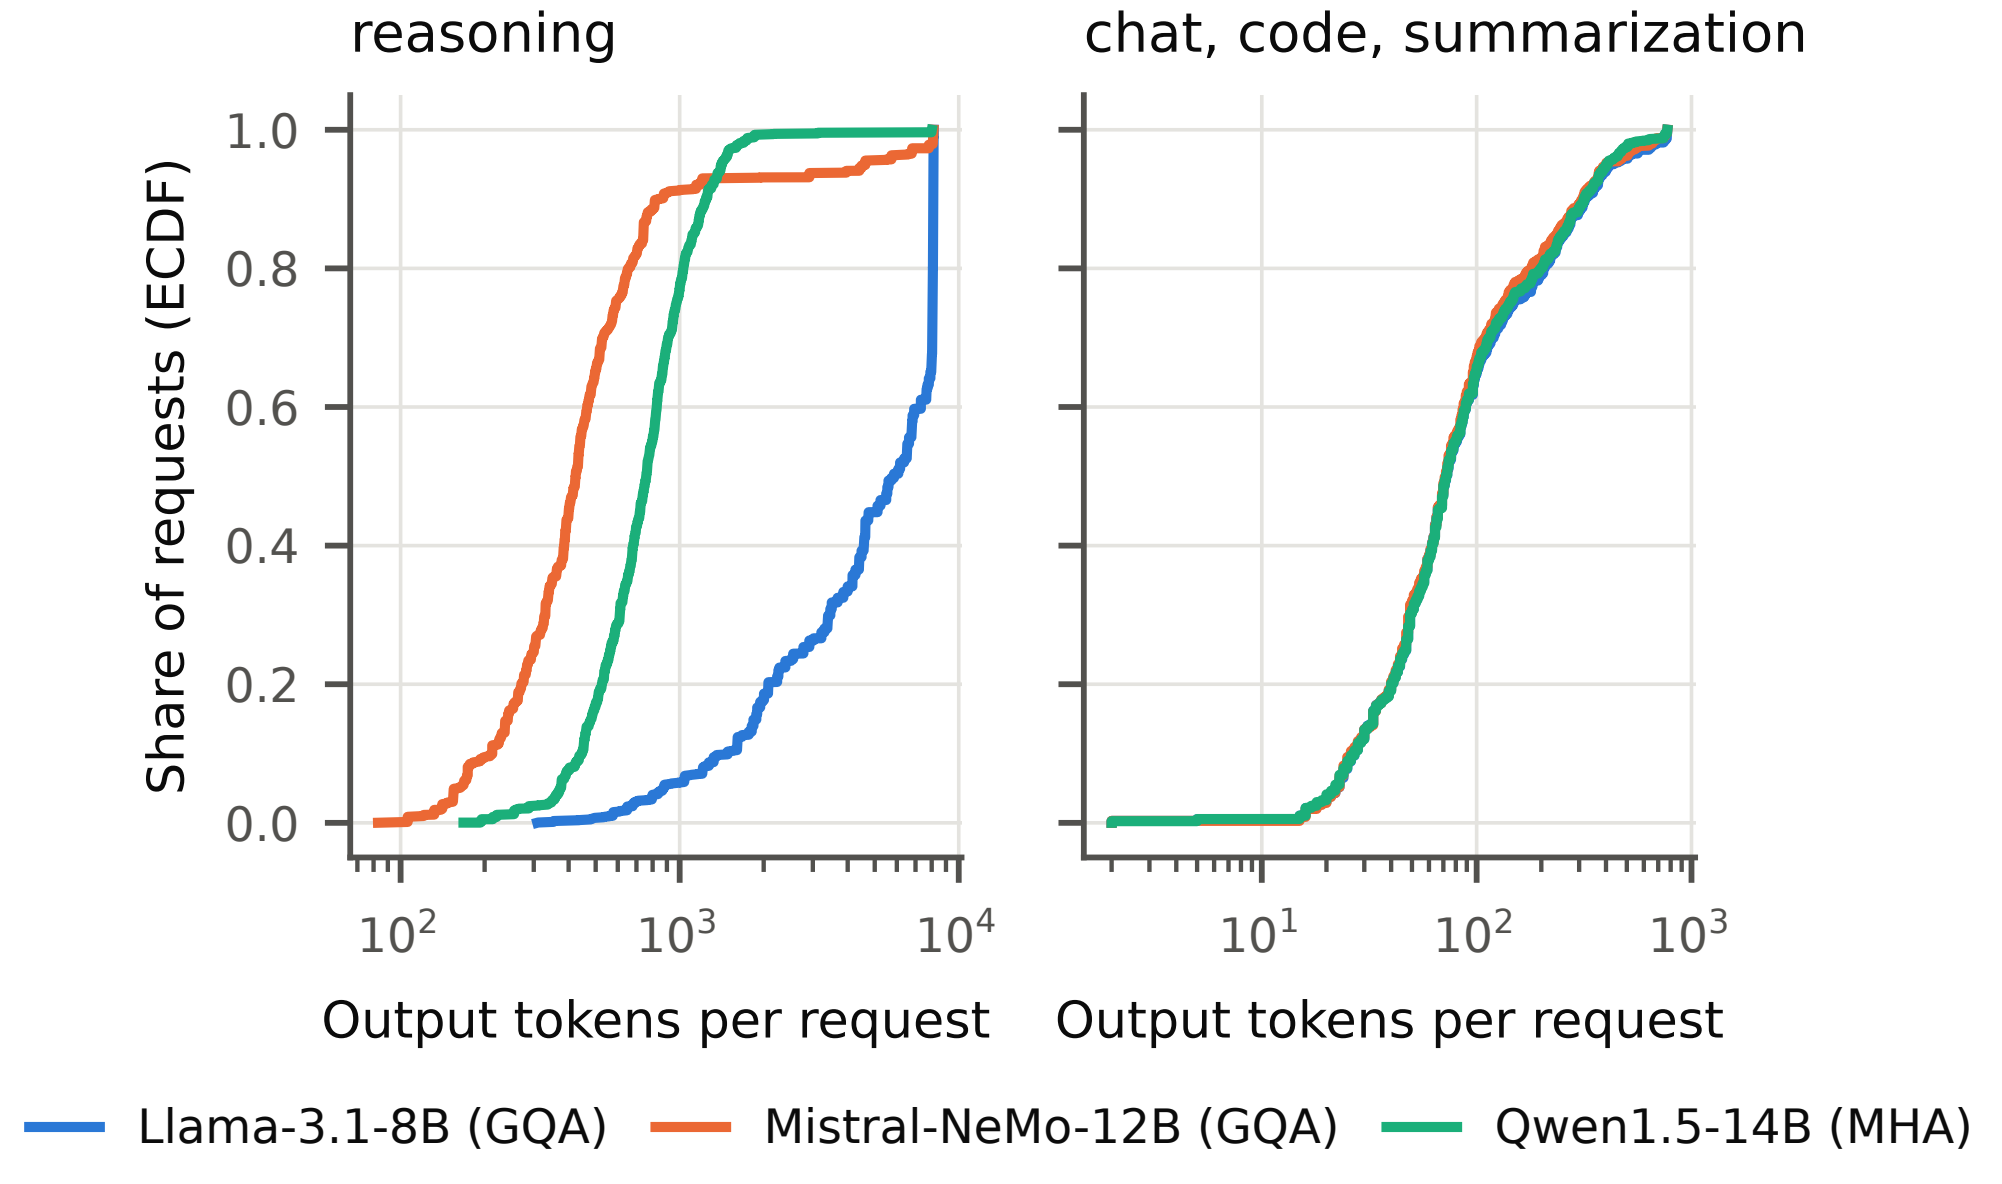
reasoning
chat, code, summarization
Share of requests (ECDF)
1.0
0.8
0.6
0.4
0.2
0.0
10^2
10^3
10^4
10^1
10^2
10^3
Output tokens per request
Output tokens per request
Llama-3.1-8B (GQA)
Mistral-NeMo-12B (GQA)
Qwen1.5-14B (MHA)