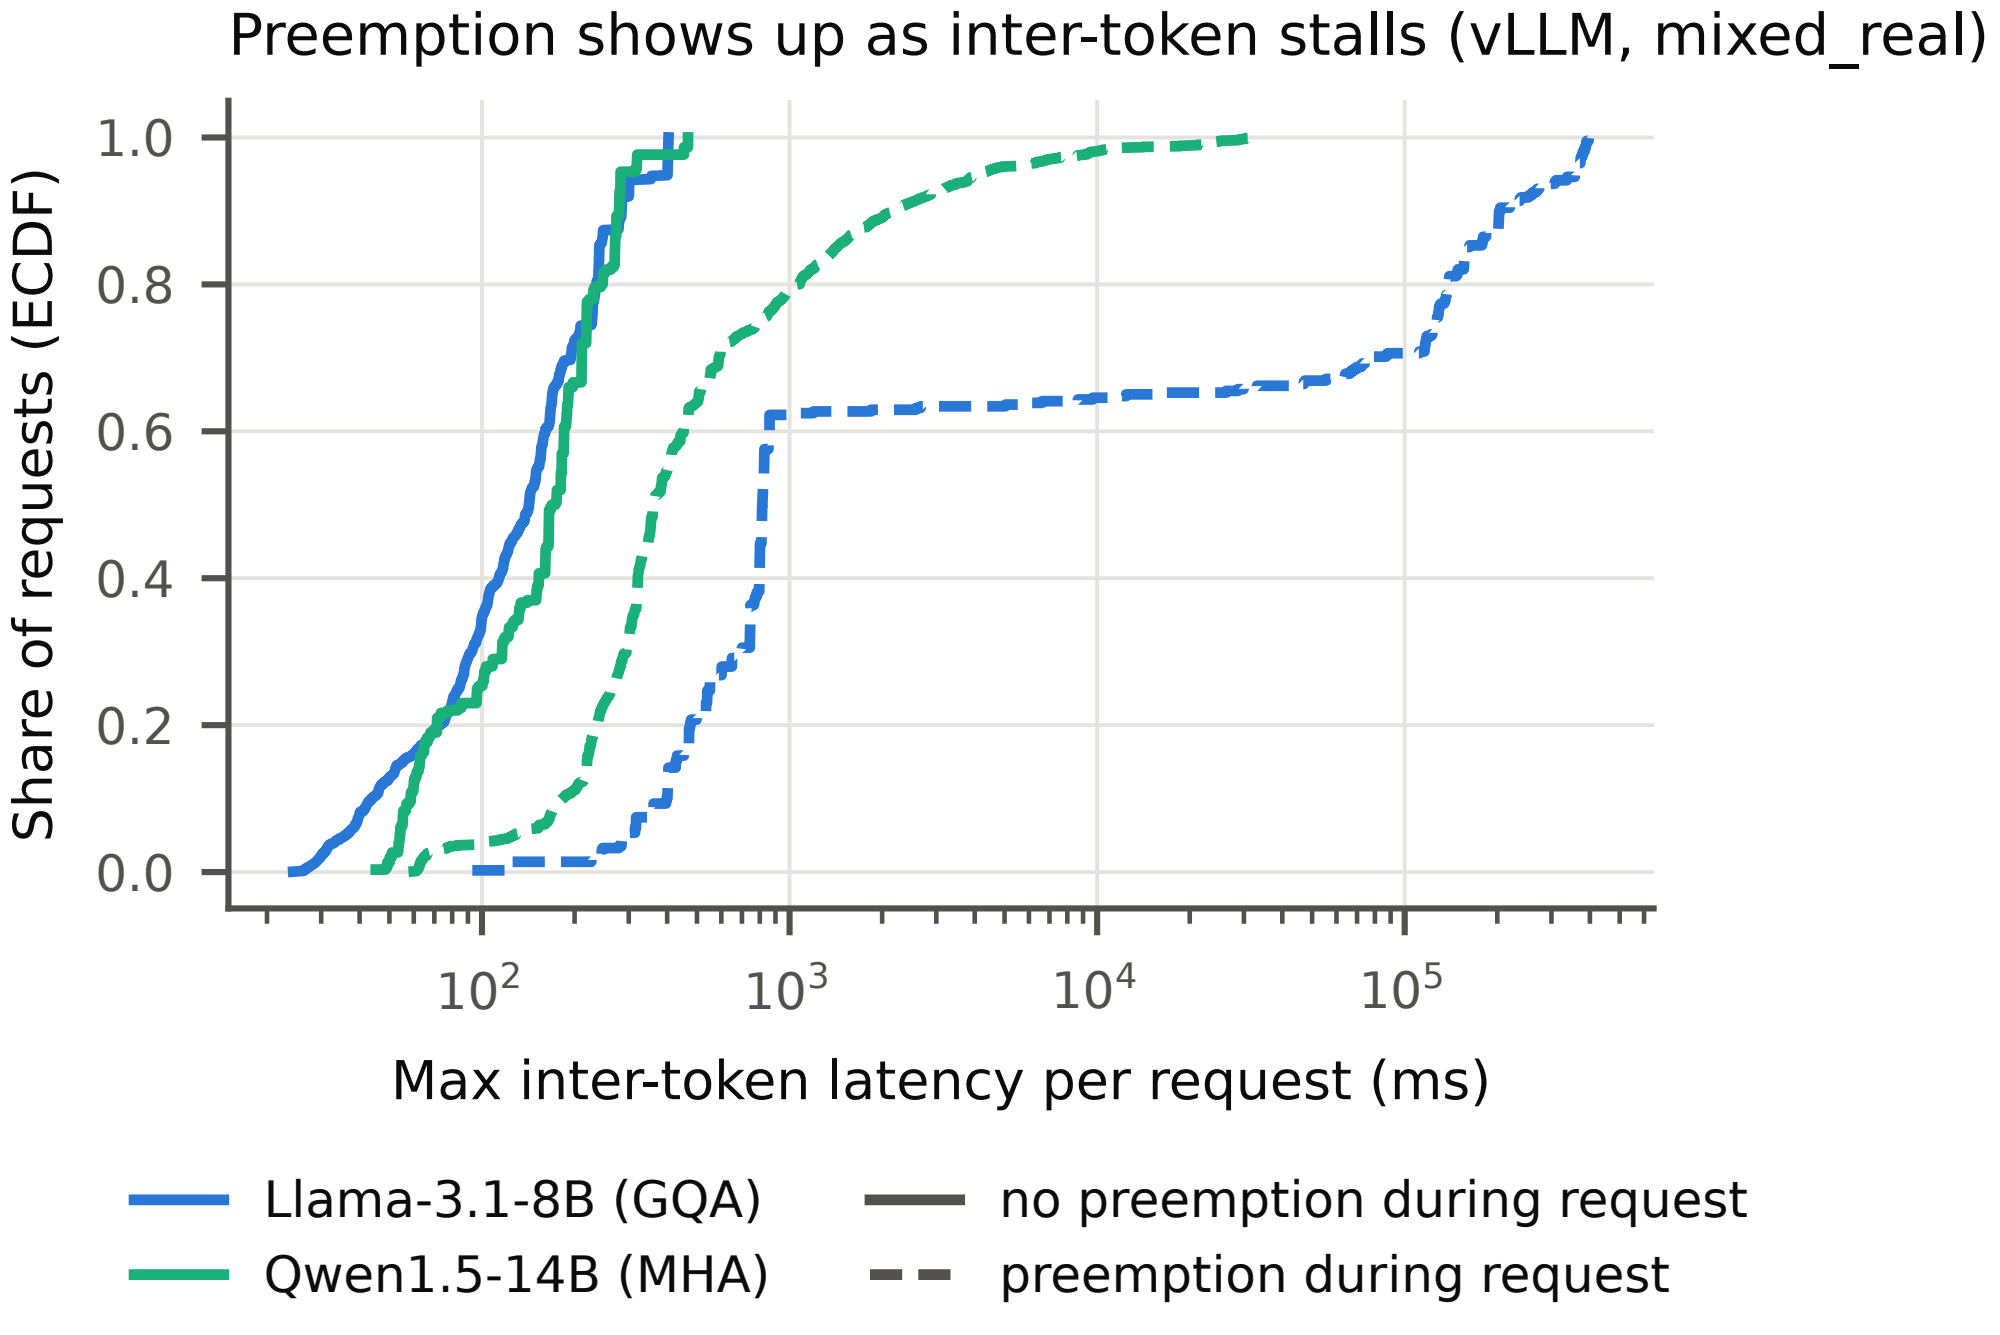
Preemption shows up as inter-token stalls (vLLM, mixed_real)
Share of requests (ECDF)
1.0
0.8
0.6
0.4
0.2
0.0
10^2
10^3
10^4
10^5
Max inter-token latency per request (ms)
Llama-3.1-8B (GQA)
Qwen1.5-14B (MHA)
no preemption during request
preemption during request