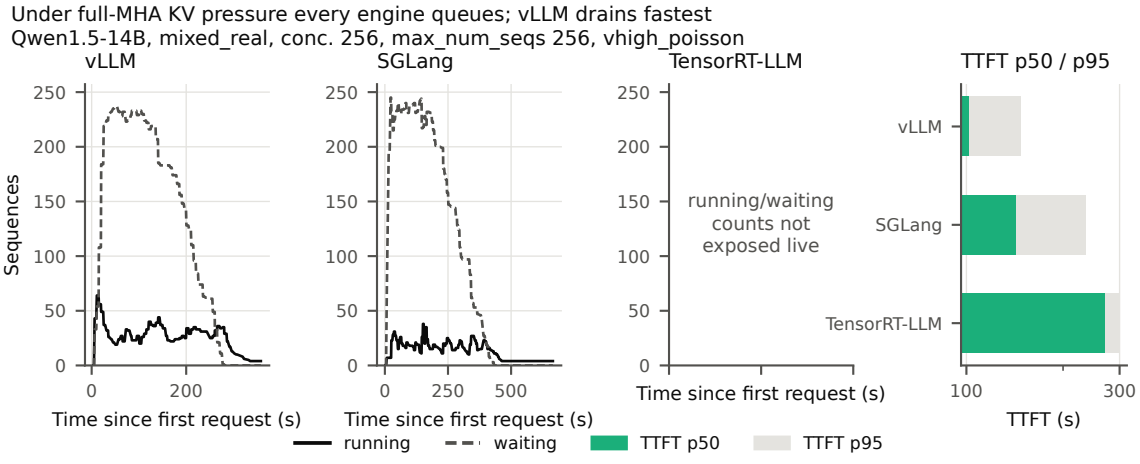
Under full-MHA KV pressure every engine queues; vLLM drains fastest
Qwen1.5-14B, mixed_real, conc. 256, max_num_seqs 256, vhigh_poisson
vLLM
SGLang
TensorRT-LLM
TTFT p50 / p95
Sequences
running/waiting counts not exposed live
Time since first request (s)
TTFT (s)
running
waiting
TTFT p50
TTFT p95