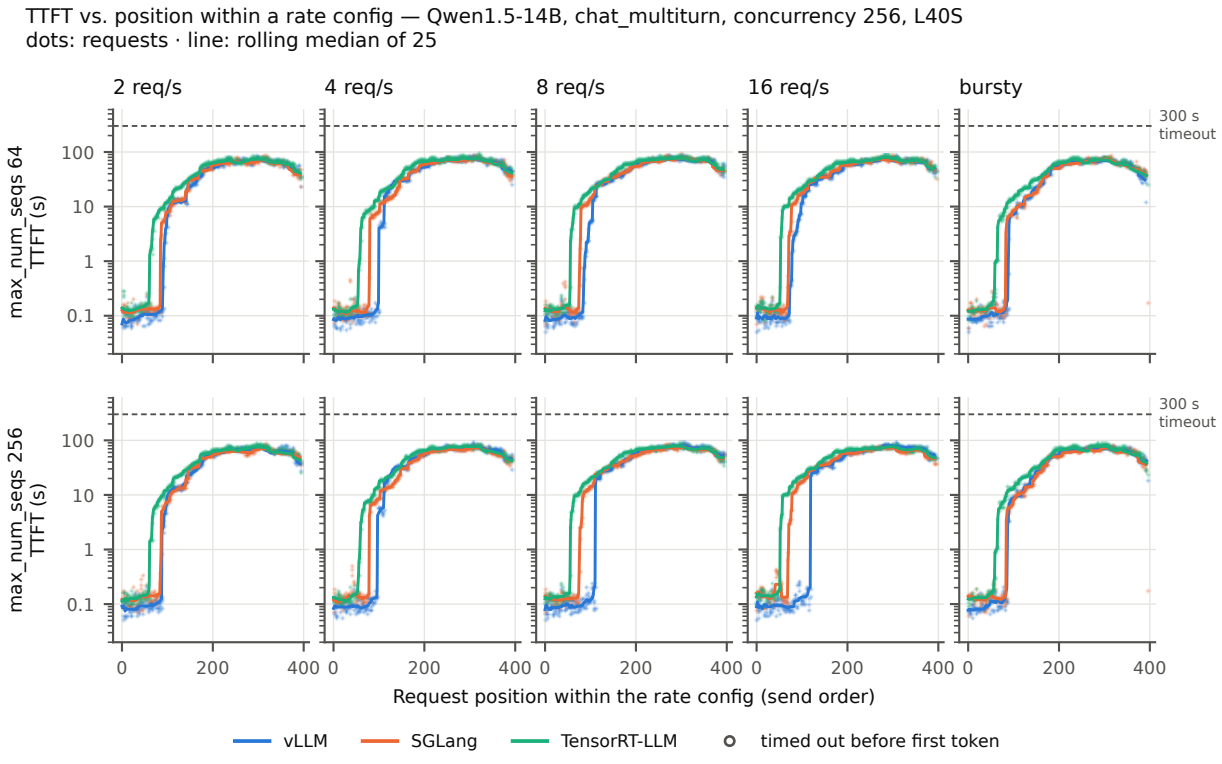

TTFT vs. position within a rate config — Qwen1.5-14B, chat_multiturn, concurrency 256, L40S
dots: requests · line: rolling median of 25
2 req/s
4 req/s
8 req/s
16 req/s
bursty
300 s timeout
max_num_seqs 64
max_num_seqs 256
TTFT (s)
100
10
1
0.1
0
200
400
Request position within the rate config (send order)
vLLM
SGLang
TensorRT-LLM
timed out before first token