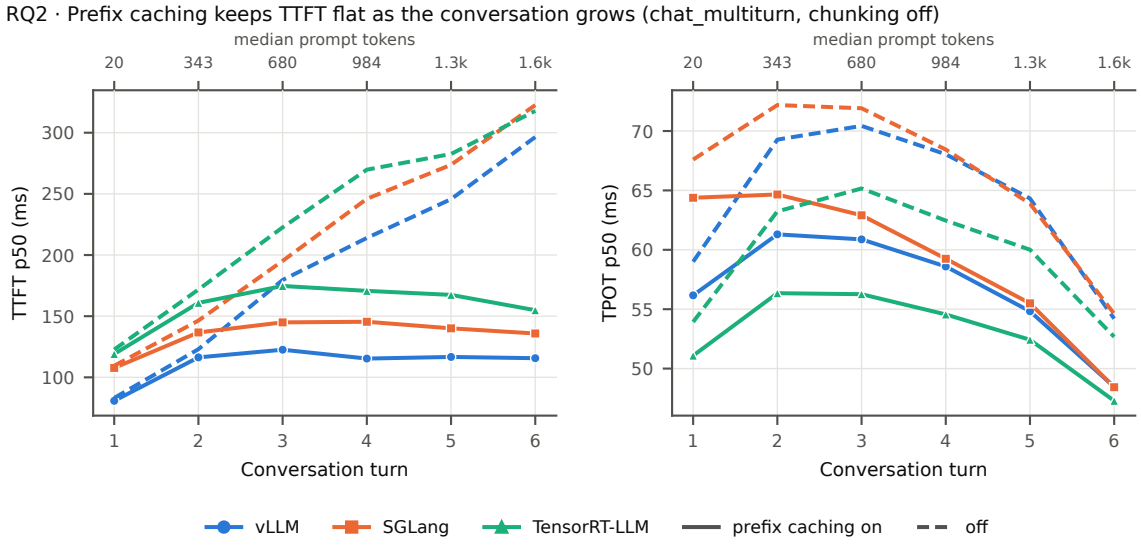
RQ2 · Prefix caching keeps TTFT flat as the conversation grows (chat_multiturn, chunking off)
median prompt tokens
20
343
680
984
1.3k
1.6k
TTFT p50 (ms)
300
250
200
150
100
1
2
3
4
5
6
Conversation turn
median prompt tokens
20
343
680
984
1.3k
1.6k
TPOT p50 (ms)
70
65
60
55
50
1
2
3
4
5
6
Conversation turn
vLLM
SGLang
TensorRT-LLM
prefix caching on
off