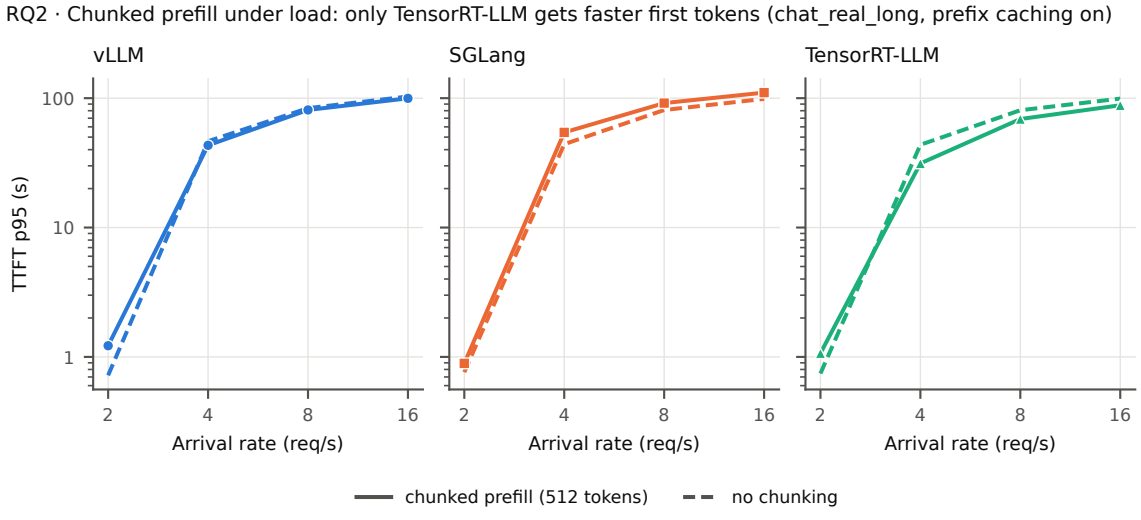
RQ2 · Chunked prefill under load: only TensorRT-LLM gets faster first tokens (chat_real_long, prefix caching on)
vLLM
SGLang
TensorRT-LLM
TTFT p95 (s)
100
10
1
2
4
8
16
Arrival rate (req/s)
chunked prefill (512 tokens)
no chunking