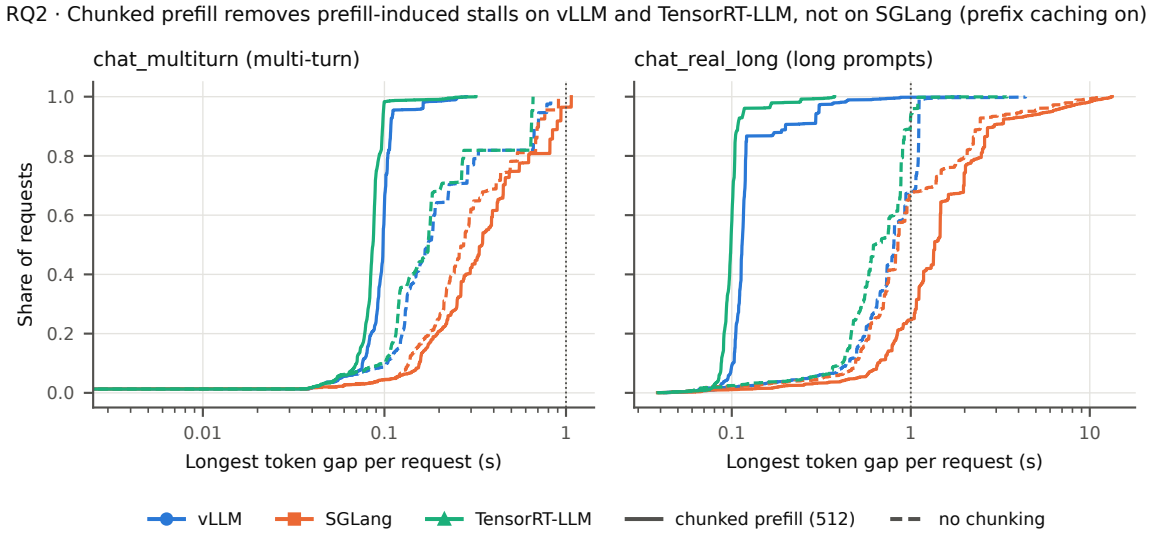
RQ2 · Chunked prefill removes prefill-induced stalls on vLLM and TensorRT-LLM, not on SGLang (prefix caching on)
chat_multiturn (multi-turn)
chat_real_long (long prompts)
Share of requests
1.0
0.8
0.6
0.4
0.2
0.0
0.01
0.1
1
10
Longest token gap per request (s)
vLLM
SGLang
TensorRT-LLM
chunked prefill (512)
no chunking