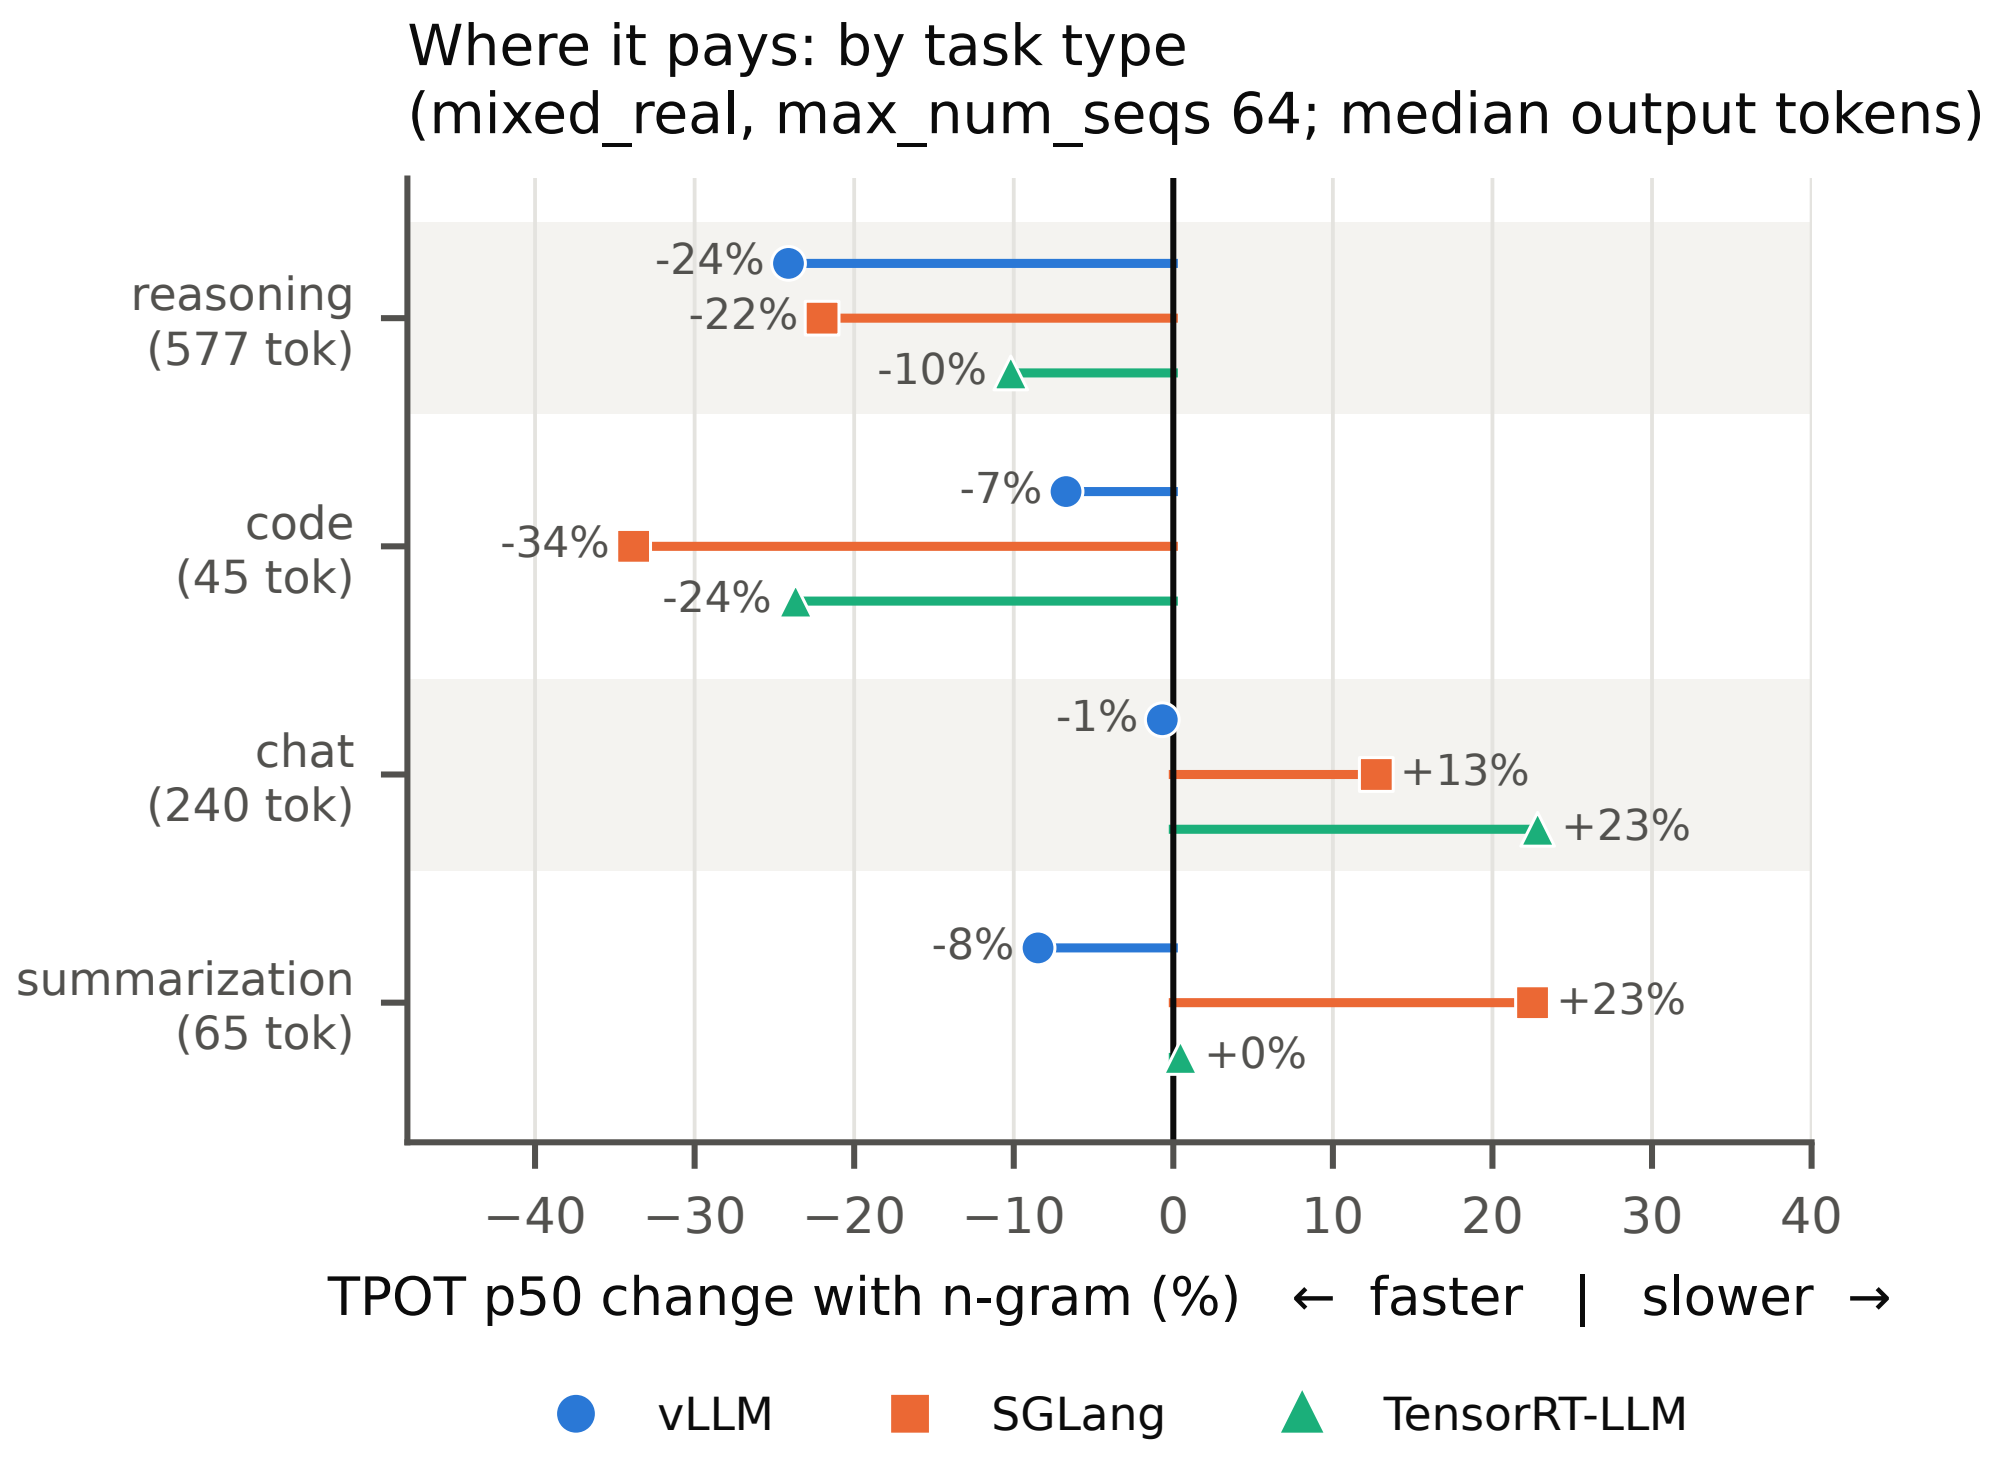
Where it pays: by task type
(mixed_real, max_num_seqs 64; median output tokens)
reasoning
(577 tok)
code
(45 tok)
chat
(240 tok)
summarization
(65 tok)
-24%
-22%
-10%
-7%
-34%
-24%
-1%
+13%
+23%
-8%
+23%
+0%
−40
−30
−20
−10
0
10
20
30
40
TPOT p50 change with n-gram (%) ← faster | slower →
vLLM
SGLang
TensorRT-LLM