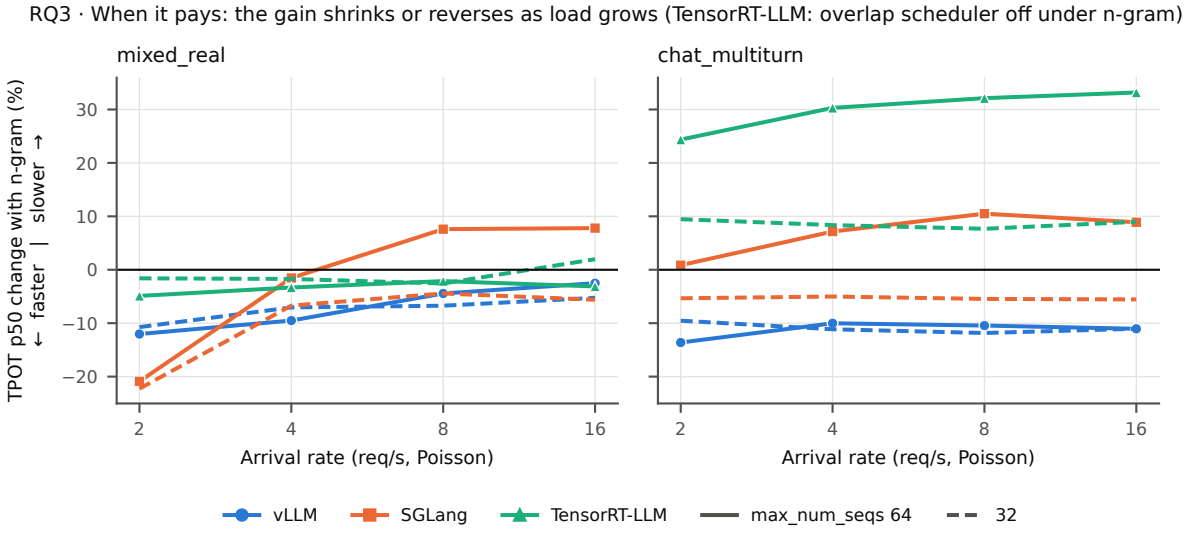
RQ3 · When it pays: the gain shrinks or reverses as load grows (TensorRT-LLM: overlap scheduler off under n-gram)
mixed_real
chat_multiturn
TPOT p50 change with n-gram (%)
← faster | slower →
30
20
10
0
−10
−20
2
4
8
16
Arrival rate (req/s, Poisson)
vLLM
SGLang
TensorRT-LLM
max_num_seqs 64
32